



ParulTM
University

NAAC A++
ACCREDITED UNIVERSITY

**FACULTY OF ENGINEERING AND TECHNOLOGY
BACHELOR OF TECHNOLOGY**

**MACHINE LEARNING
(303105354)**

**VI SEMESTER
6B3**

COMPUTER SCIENCE AND ENGINEERING DEPARTMENT

Laboratory Manual

CERTIFICATE

This is to certify that Mr./Ms. XXXXXXXXXXXXXX with
Enrollment No. 21030XXXXXXXXX has successfully completed her laboratory
experiments in **Machine Learning Laboratory (303105354)** From the
Department of Computer Science and Engineering during the
academic year **2025 – 2026**



Date of Submission:

Staff In Charge:

Head of Department:

TABLE OF CONTENT

SR. NO	PRACTICAL LIST	PAGE NO.	START DATE	END DATE	SIGN	MARKS
1	Dealing with Data using NumPy, Pandas, Statistics library.					
2	Data Analysis & Visualization on Diwali Sales Dataset.					
3	Implementation on <ul style="list-style-type: none"> a) Simple Linear Regression b) Multi-Linear Regression c) Logistic Regression 					
4	Implement the naïve Bayesian classifier for a sample training data set stored as a .CSV file. Compute the accuracy of the classifier, considering a few test data sets.					
5	Assuming a set of documents that need to be classified, use the naïve Bayesian Classifier model to perform this task.					
6	Decision tree-based ID3 algorithm.					
7	Write a program to implement the K-Nearest Neighbor algorithm to classify the iris data set.					
8	Apply EM algorithm to cluster a set of data stored in a .CSV file. Use the same data set for clustering using k-Means algorithm.					
9	Write a program to construct a Bayesian network considering medical data. Use this model to demonstrate the diagnosis of heart patients using standard Heart Disease Data set.					
10	Compare the various supervised learning algorithm by using appropriate dataset.					
11	Compare the various Unsupervised learning algorithm by using the appropriate datasets.					
12	Build an Artificial Neural Network by implementing the Backpropagation algorithm and test the same using appropriate data sets.					

Experiment - 1

Aim: Dealing With Data Using Stats Numpy and Pandas Library.

Code:

```
import statistics as st
a=[20, 10, 50, 10, 21, 90]
s=st.mode(a)
print("Mode:",s)
```

```
⇒ Mode: 10
```

```
import statistics as st
a=[20, 10, 50, 10, 21, 90]
m=st.mean(a)
print("Mean:",m)
```

```
⇒ Mean: 33.5
```

```
pip install numpy
```

```
⇒ Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.25.2)
```

```
import numpy as np
b=np.array([[10, 20, 30, 50, 10],[1, 2, 3, 5, 6]])
print(b)
print("Dim:",b.ndim)
```

```
⇒ [[10 20 30 50
     10]]
```

```
[ 1  2  3  5  6]
Dim: 2
```

```
import numpy as np
b=np.array([[10, 20, 30, 50, 10],[1, 2, 3, 5, 6],[1,1,1,1,1]])
print(b)
print("Dim:",b.ndim)
print("Shape:",b.shape)
print("Size:",b.size)
print("Item Size:",b.itemsize)
print("Data type:",b.dtype)
```

```
⇒ [[10 20 30 50 10]
     [ 1  2  3  5  6]
     [ 1  1  1  1  1]]
Dim: 2
Shape: (3, 5)
Size: 15
Item Size: 8
Data type: int64
```

```
c=np.empty((5,5))
```

```

array([[4.85494045e-310, 0.00000000e+000, 6.28083585e-038,
       2.16385130e+190, 1.94918964e-153],
      [4.90838520e+252, 1.30304358e-142, 1.06396443e+224,
       1.70100503e+256, 5.49109388e-143],
      [9.30537467e+199, 2.57018365e-057, 1.56067931e+184, 3.53933111e-
       061, 4.85198335e-033],
      [4.40779861e-143, 5.49810101e-143, 6.19651793e-091,
       1.49655354e-076, 1.57712453e-052],
      [4.29600452e-096, 4.82337433e+228, 3.67152456e-062, 6.01346930e-
       154, 1.81816158e-321]])

d=np.arange(0,100,2
) d

array([ 0,  2,  4,  6,  8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32,
       34, 36, 38, 40, 42, 44, 46, 48, 50, 52, 54, 56, 58, 60, 62, 64, 66,
       68, 70, 72, 74, 76, 78, 80, 82, 84, 86, 88, 90, 92, 94, 96, 98])

e=np.reshape(d,(10,5))

f=np.reshape(e,(2,5,5))

print(f)

print("Last element:",d[10])

print("Element upto 10:",d[:10])

array Element upto 10: [ 0  2  4  6  8 10 12 14 16 18]

print("Last 5 element:",d[-5:])

array Last 5 element: [90 92 94 96 98]

print("Alternative Numbers:",d[::-2])

array Alternative Numbers: [ 0  4  8 12 16 20 24 28 32 36 40 44 48 52 56 60 64 68 72 76 80 84 88 92 96]

pip install pandas

Collecting
panda
  Downloading panda-0.3.1.tar.gz
    (5.8 kB) Preparing metadata
      (setup.py) ... done
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from panda)
(67.7.2)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from panda)
(2.31.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages
(from requests->panda) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.10/dist-packages (from requests->panda) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from

```

```
requests->panda) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from
requests->panda) (2024.6.2)Building wheels for collected packages: panda
Building wheel for panda (setup.py) ... done
Created wheel for panda: filename=panda-0.3.1-py3-none-any.whl size=7239
sha256=403d903fdbfc2160794713c5b6c1673297582399c068021f Stored in directory:
/root/.cache/pip/wheels/0e/8b/c3/ff9cbde1ffd8071cff8367a86f0350a1ce30a8d31b6a432e9
Successfully built panda
Installing collected
packages: pandaSuccessfully
<   ██████████   >
```

installed panda-0.3.1

```
import pandas as pd
p_dict={'pid':[1,2,3,4,5],
'pid':[1,2,3,4,5],
'pname':['A','B','C','D','E'], 'price':[10,20,30,40,50],
'vendor':['V1','V2','V3','V4','V5']}
p_dict
```

```
→  {'pid': [1, 2, 3, 4, 5],
  'pname': ['A', 'B', 'C', 'D', 'E'], 'price': [10, 20, 30, 40, 50],
  'vendor': ['V1', 'V2', 'V3', 'V4', 'V5']}
```

```
grt=pd.DataFrame(p_dict) grt
```

```
→  pid  pname  price  vendor
  0    1      A     10      V1
  1    2      B     20      V2
  2    3      C     30      V3
  3    4      D     40      V4
```

Experiment - 2

Aim: Data Analysis and Visualisation on Diwali Sales Database.

Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df=pd.read_csv('Diwali Sales
Data.csv',encoding='unicode_escape') df
```

	User_ID	Cust_name	Product_ID	Gender	Age	Marital_Status	State	Zone	Occupation	Product_Category	0
					Group	Product_Catagory					
0	1002903	Sanskriti	P00125942	F	26-35	28	0	Maharashtra	Western Healthcare	Auto	
1	1000732	Kartik	P00110942	F	26-35	35	1	Andhra Pradesh	Southern Govt	Auto	
2	1001990	Bindu	P00118542	F	26-35	35	1	Uttar Pradesh	Central Automobile	Auto	
3	1001425	Sudevi	P00237842	M	0-17	16	0	Karnataka	Southern Construction	Auto	
4	1000588	Joni	P00057942	M	26-35	28	1	Gujarat	Western Food Processing	Auto	
...
11246	1000695	Manning	P00296942	M	18-25	19	1	Maharashtra	Western Chemical	Office	
11247	1004089	Reichenbach	P00171342	M	26-35	33	0	Haryana	Northern Healthcare	Veterinary	
11248	1001209	Oshin	P00201342	F	36-45	40	0	Madhya Pradesh	Central Textile	Office	
11249	1004023	Noonan	P00059442	M	36-45	37	0	Karnataka	Southern Agriculture	Office	
11250	1002744	Brumley	P00281742	F	18-25	19	0	Maharashtra	Western Healthcare	Office	

11251 rows × 15 columns

Next steps:

[Generate code with df](#)

[View recommended plots](#)

df.shape

(11251, 15)

df.head()

	User_ID	Cust_name	Product_ID	Gender	Age	Marital_Status	State	Zone	Occupation	Product_Category	Orders
					Group	Age					
0	1002903	Sanskriti	P00125942	F	26-35	28	0	Maharashtra	Western Healthcare	Auto	1
1	1000732	Kartik	P00110942	F	26-35	35	1	Andhra Pradesh	Southern Govt	Auto	3
2	1001990	Bindu	P00118542	F	26-35	35	1	Uttar Pradesh	Central Automobile	Auto	3
3	1001425	Sudevi	P00237842	M	0-17	16	0	Karnataka	Southern Construction	Auto	2

4 1000588 Joni P00057942 M 26-35 28 1 Gujarat Western Food Processing Auto 2

Next steps: [Generate code with df](#) [View recommended plots](#)

df.info()

```
→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 11251 entries, 0 to 11250
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   User_ID     11251 non-null   int64  
 1   Cust_name   11251 non-null   object  
 2   Product_ID  11251 non-null   object  
 3   Gender      11251 non-null   object  
 4   Age Group   11251 non-null   object  
 5   Age          11251 non-null   int64  
 6   Marital_Status 11251 non-null   int64  
 7   State        11251 non-null   object  
 8   Zone         11251 non-null   object  
 9   Occupation   11251 non-null   object  
 10  Product_Category 11251 non-null   object  
 11  Orders       11251 non-null   int64  
 12  Amount       11239 non-null   float64 
 13  Status       0 non-null      float64 
```

```
14 unnamed1
0 non-null
float64 dtypes: float64(3), int64(4),
object(8)
memory usage: 1.3+ MB
```

```
df.drop(['Status', 'unnamed1'], axis=1, inplace=True)
df.drop(['Status', 'unnamed1'], axis=1, inplace=True)
```

```
pd.isnull(df).sum()
```

```
→ User_ID      0
Cust_name    0
Product_ID   0
Gender       0
Age Group    0
Age          0
Marital_Status 0
State         0
Zone          0
Occupation   0
Product_Category 0
Orders        0
Amount        12
dtype: int64 
```

```
df.dropna(inplace=True)
```

```
df.info()
```

```
→ <class 'pandas.core.frame.DataFrame'> Index: 11239 entries, 0 to 11250
Data columns (total 13 columns):
```

```
#           Column Non-Null Count Dtype 
----- 
0          User_ID    11239 non-null   int64 
1         Cust_name   11239 non-null   object 
2        Product_ID   11239 non-null   object 
3          Gender     11239 non-null   object 
4      Age Group    11239 non-null   object 
5            Age     11239 non-null   int64 
6      Marital_Status 11239 non-null   int64 
7            State    11239 non-null   object 
8            Zone     11239 non-null   object 
9      Occupation    11239 non-null   object 
10     Product_Category 11239 non-null   object 
11          Orders    11239 non-null   int64 
12          Amount    11239 non-null   float64 

dtypes: float64(1), int64(4), 
object(8) memory usage: 1.2+ MB
```

```
df['Amount']=df['Amount'].astype('int')
```

```
df['Amount'].d
```

```
types ↴
```

```
dtype('int64')
```

```
df.columns
```

```
Index(['User_ID', 'Cust_name', 'Product_ID', 'Gender', 
       'Age Group', 'Age', 'Marital_Status', 'State', 
       'Zone', 'Occupation', 'Product_Category', 'Orders', 
       'Amount'], 
      dtype='object')
```

```
df.info()
```

```
→ <class 'pandas.core.frame.DataFrame'> Index: 11239 entries, 0 to 11250 
Data columns (total 13 columns):
```

```
#           Column Non-Null Count Dtype 
----- 
0          User_ID    11239 non-null   int64 
1         Cust_name   11239 non-null   object 
2        Product_ID   11239 non-null   object 
3          Gender     11239 non-null   object 
4      Age Group    11239 non-null   object 
5            Age     11239 non-null   int64 
6      Marital_Status 11239 non-null   int64 
7            State    11239 non-null   object 

8    Zone        11239 non-null   object 
9  Occupation   11239 non-null   object 
10 Product_Category 11239 non-null   object 
11  Orders      11239 non-null   int64 
12  Amount      11239 non-null   float64 

dtypes: int64(5), object(8) 
memory usage: 1.2+ MB
```

df.describe()

	User_ID	Age	Marital_Status	Orders	Amount
count	1.123900e+04	11239.000000	11239.000000	11239.000000	11239.000000
mean	1.003004e+06	35.410357	0.420055	2.489634	9453.610553
std	1.716039e+03	12.753866	0.493589	1.114967	5222.355168
min	1.000001e+06	12.000000	0.000000	1.000000	188.000000
25%	1.001492e+06	27.000000	0.000000	2.000000	5443.000000
50%	1.003064e+06	33.000000	0.000000	2.000000	8109.000000
75%	1.004426e+06	43.000000	1.000000	3.000000	12675.000000
max	1.006040e+06	92.000000	1.000000	4.000000	23952.000000

df.describe(include='all')

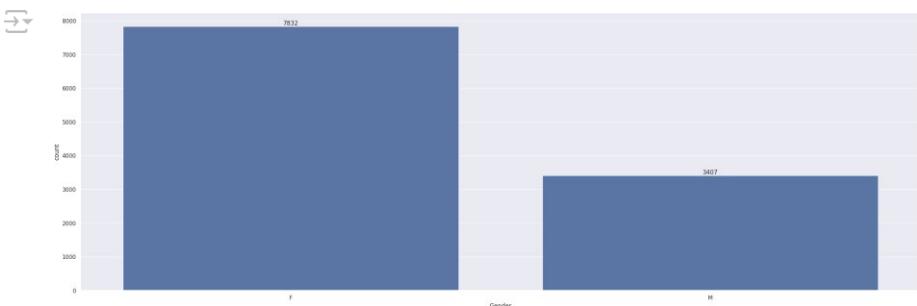
count	1.123900e+04	11239	11239	11239	11239	11239.000000
unique	NaN	1250	2350	2	7	NaN
top	NaN	Vishakha	P00265242		F	26-35
freq	NaN	42	53	7832	4541	NaN
	mean	1.003004e+06		NaN	NaN	NaN
35.410357		0.4200				
std	1.716039e+03	NaN	NaN	NaN	12.753866	0.4935
min	1.000001e+06	NaN	NaN	NaN	12.000000	0.0000
25%	1.001492e+06	NaN	NaN	NaN	27.000000	0.0000
50%	1.003064e+06	NaN	NaN	NaN	33.000000	0.0000
75%	1.004426e+06	NaN	NaN	NaN	43.000000	1.0000
max	1.006040e+06	NaN	NaN	NaN	92.000000	1.0000

df[['Age', 'Orders', 'Amount']].describe()

	Age	Orders	Amount
count	11239.000000	11239.000000	11239.000000
mean	35.410357	2.489634	9453.610553
std	12.753866	1.114967	5222.355168
min	12.000000	1.000000	188.000000
25%	27.000000	2.000000	5443.000000
50%	33.000000	2.000000	8109.000000
75%	43.000000	3.000000	12675.000000
max	92.000000	4.000000	23952.000000

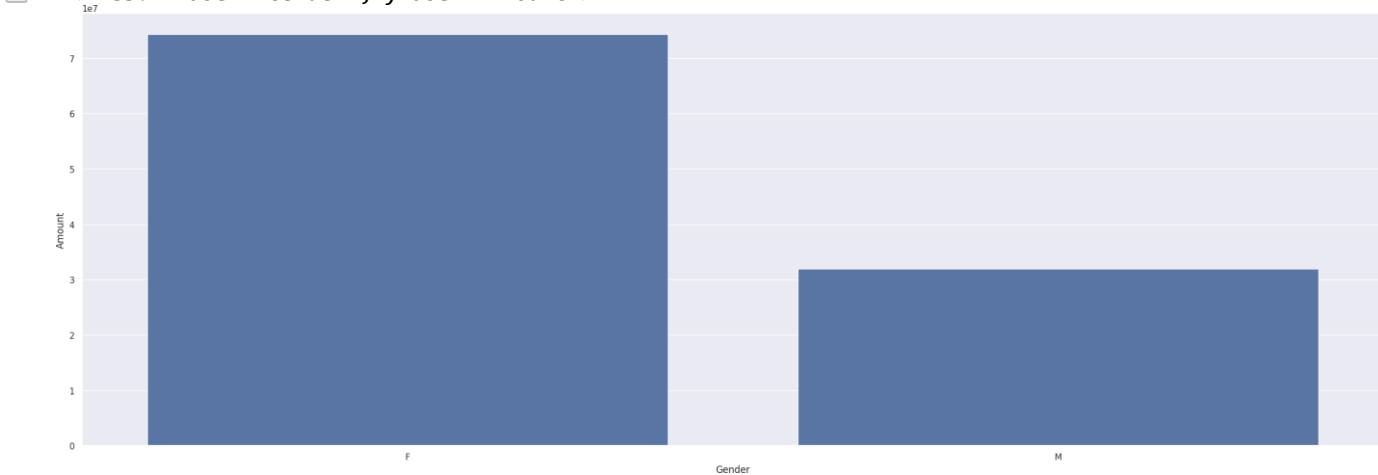
#plotting a bar chart for Gnder and it's Count ax=sns.countplot(x='Gender', data=df)

```
for bars in ax.containers: ax.bar_label(bars)
```

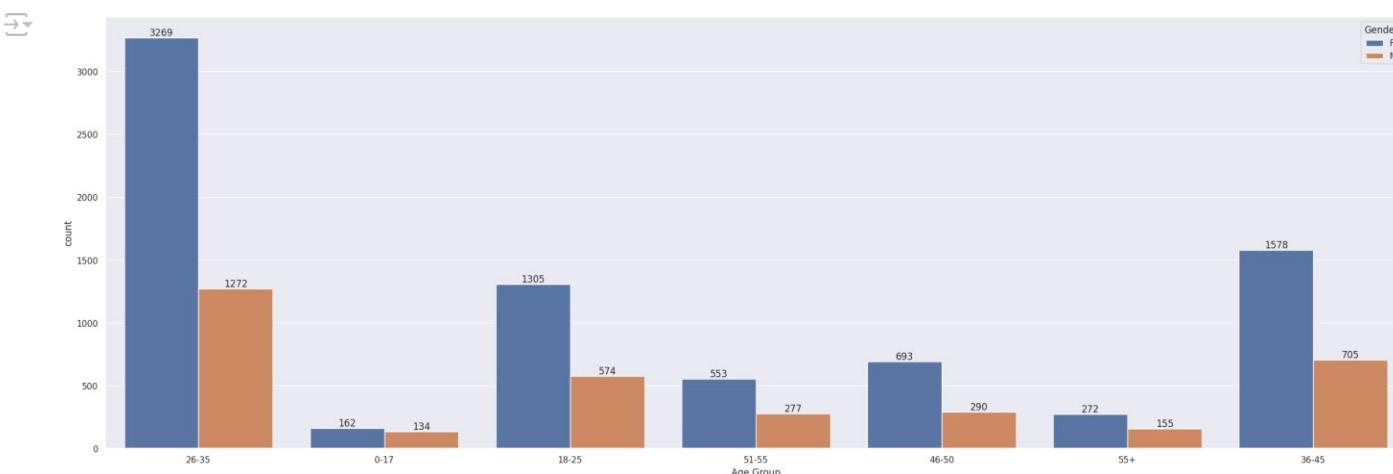


```
sales_gen=df.groupby(['Gender'],as_index=False)[['Amount']].sum().sort_values(by='Amount',ascending=False)sns.barplot(x='Gender',y='Amount',data=sales_gen)
```

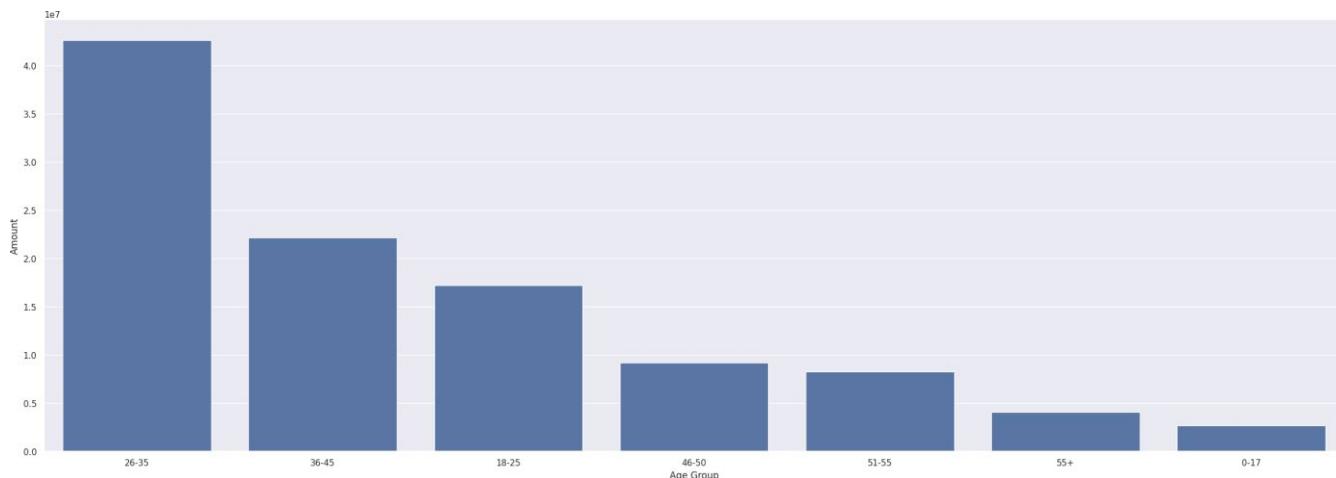
```
<Axes: xlabel='Gender', ylabel='Amount'>
```



```
ax=sns.countplot(data=df,x='Age Group',hue='Gender')  
for bars in ax.containers:  
    ax.bar_label(bars)
```

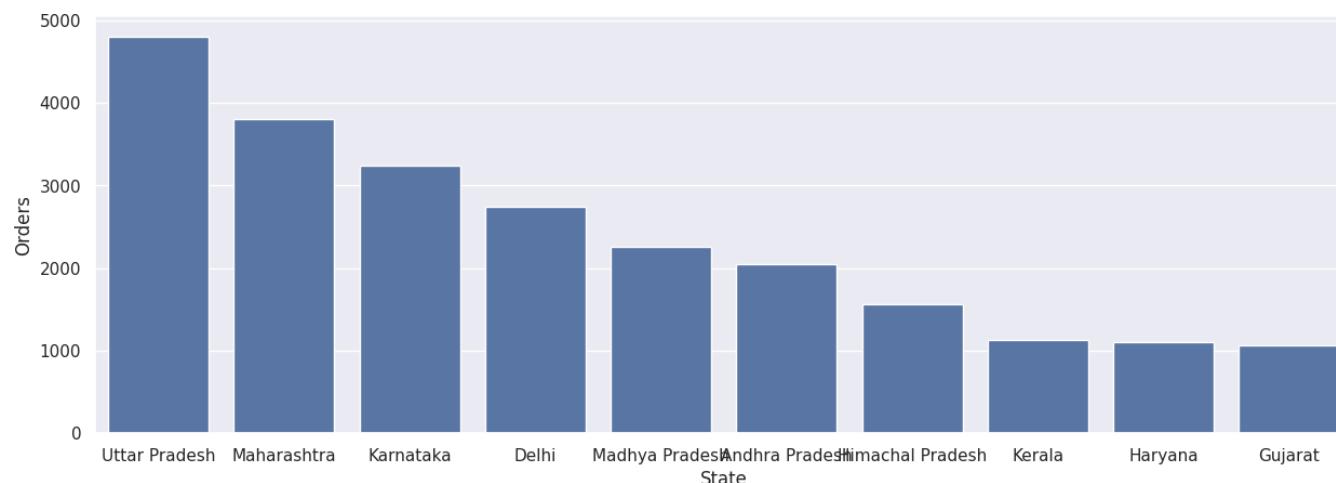


```
sales_age=df.groupby(['Age Group'],as_index=False)[['Amount']].sum().sort_values(by='Amount',ascending=False)sns.barplot(x='Age Group',y='Amount',data=sales_age)
```



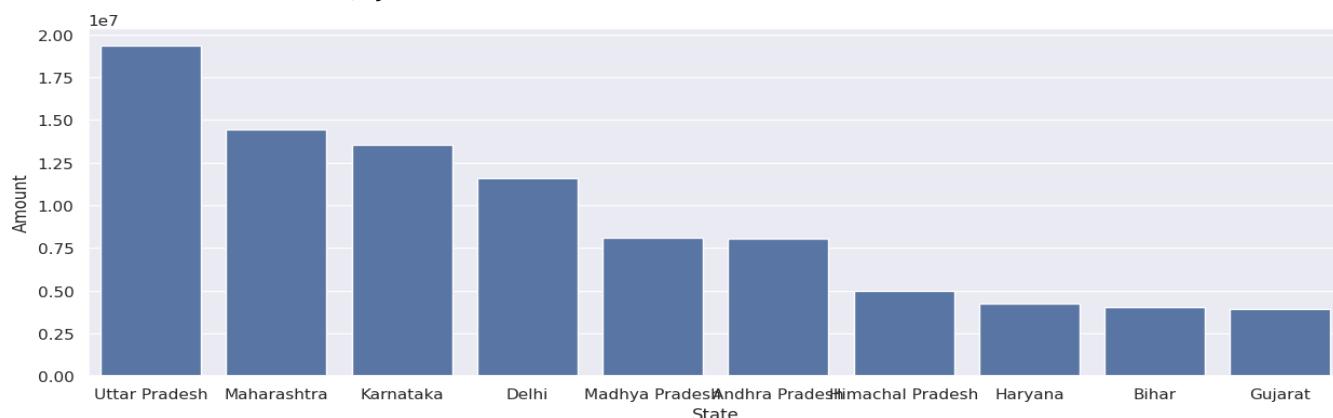
```
<Axes: xlabel='Age Group', ylabel='Amount'
sales_states=df.groupby(['State'],as_index=False)[['Orders']].sum().sort_values(by='Orders', ascending=False).head(10)sns.set(rc={'figure.figsize':(15,5)})
sns.barplot(x='State',y='Orders',data=sales_states)
```

```
<Axes:
xlabel='State',ylabel='Orders'>
```



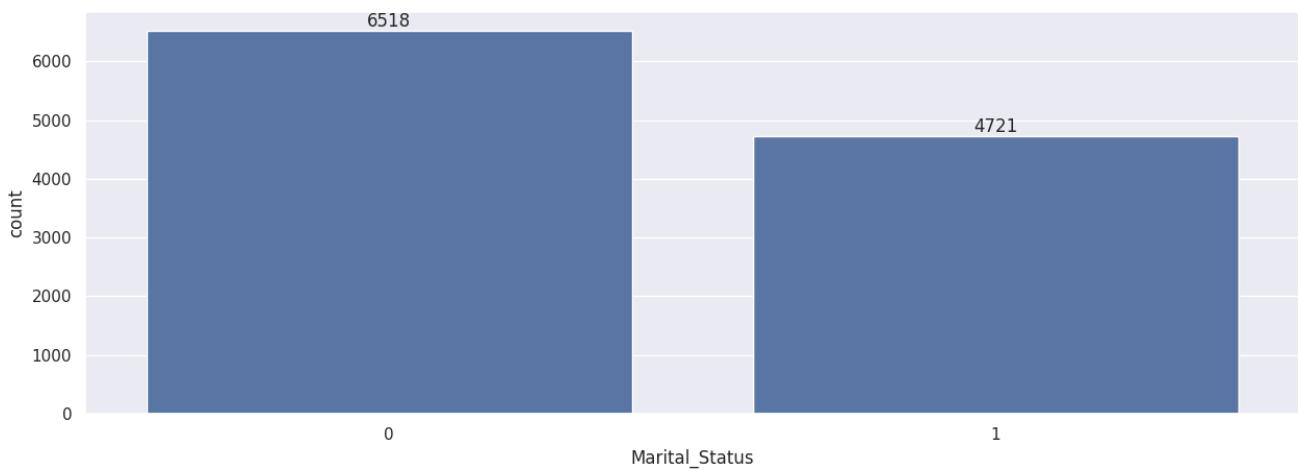
```
sales_state=df.groupby(['State'],as_index=False)[['Amount']].sum().sort_values(by='Amount', ascending=False).head(10)sns.set(rc={'figure.figsize':(15,5)})
sns.barplot(x='State',y='Amount',data=sales_state)
```

```
<Axes: xlabel='State', ylabel='Amount'>
```



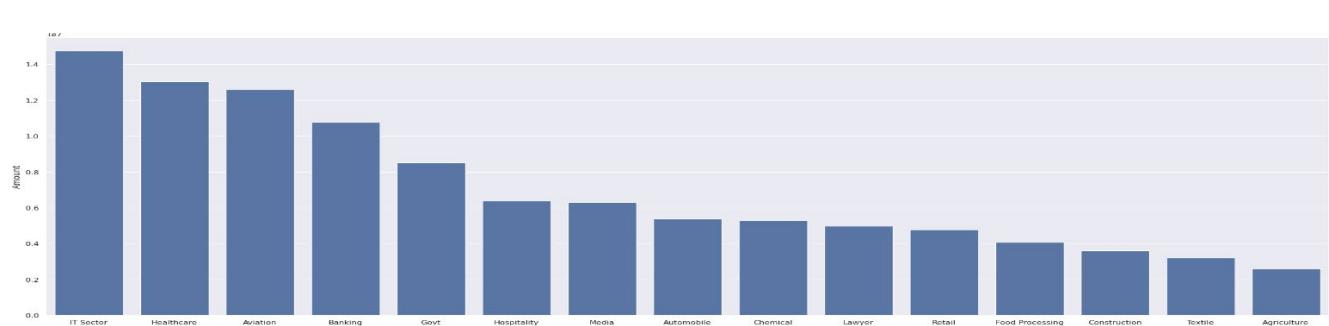
```
ax=sns.countplot(data=df,x='Marital_Status')
sns.set(rc={'figure.figsize':(7,5)})
```

```
for bars in ax.containers:  
    ax.bar_label(bars)
```



```
sns.set(rc={'figure.figsize':(30,10)})  
sales_state=df.groupby(['Occupation'],as_index=False)[['Amount']].sum().sort_values(by='Amount',ascending=False)sns.barplot(x='Occupation',y='Amount',data=sales_state)
```

```
<Axes: xlabel='Occupation', ylabel='Amount'>
```



Experiment - 3

Aim: Implementation On Linear Regression And Logistic Regression.

Code:

Liner Regression:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as
plt
import seaborn as sns

df=pd.read_csv('placement.csv',encoding='unico
de_escape')
df
```

	cgpa	packag
0	6.89	3.26
1	5.12	1.98
2	7.82	3.25
3	7.42	3.67
4	6.94	3.57
...
195	6.93	2.46
196	5.89	2.57
197	7.21	3.24
198	7.63	3.96
199	6.22	2.33

200 rows x 2 columns

[Generate code with df](#)

[View recommended plots](#)

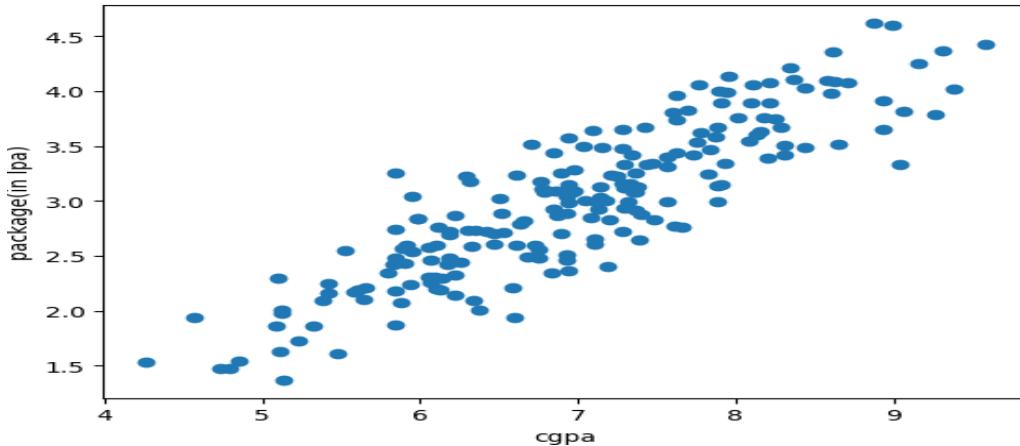
Next steps:

```
df.info()

class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 2 columns):
 #   Column Non-Null Count Dtype
---- 
 0   cgpa    200 non-null float64
 1   package  200 non-null float64
               dtypes: float64(2)
memory usage: 3.2 KB
```

```
plt.scatter(df['cgpa'],df['package'])
plt.xlabel('cgpa')
plt.ylabel('package(in lpa)')
```

Text(0, 0.5, 'package(in lpa)')



```
x=df.iloc[:,0:1]
y=df.iloc[:,-
1] y
```

Text(0 3.26
1 1.98

```
2      3.25
3      3.67
4      3.57
195     ...
196     2.46
197     2.57
198     3.24
199     3.96
200    2.33
```

Name: package, Length: 200, dtype: float64

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=2)

from sklearn.linear_model import
LinearRegression lr=LinearRegression()
lr.fit(x_train,y_train)
```

Text(LinearRegression()
LinearRegression()

x_test

Text(

115	8.35
35	6.87
12	8.94

92	7.90
13	6.93
126	5.91
174	7.32
2	7.82
44	5.09
3	7.42
113	6.94
14	7.73
23	6.19
25	7.28
6	6.73
134	7.20
165	8.21
173	6.75
45	7.87
65	7.60
48	8.63
122	5.12
178	8.15
64	7.36
9	8.31
57	6.60
78	6.59
71	7.47
128	7.93
176	6.29
131	6.37
53	6.47

Next steps:

Generate code with
x_test

View recommended
plots

y_test

112	4.10
29	3.49
182	2.08
199	2.33
193	1.94
85	1.48
10	1.86

35 2.87

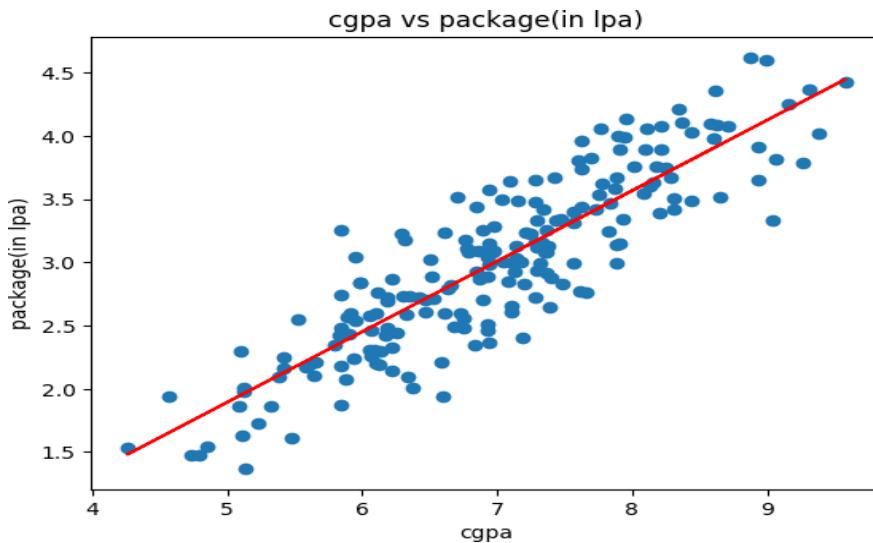
```
12  3.65
92  4.00
13  2.89
126 2.60
174 2.99
  2 3.25
 44 1.86
   3 3.67
113 2.37
 14 3.42
 23 2.48
 25 3.65
   6 2.60
134 2.83
165 4.08
173 2.56
 45 3.58
 65 3.81
 48 4.09
122 2.01
178 3.63
 64 2.92
   9 3.51
 57 1.94
 78 2.21
 71 3.34
128 3.34
176 3.23
131 2.01
 53 2.61
Name: package, dtype: float64
```

```
lr.predict(x_test.iloc[0].values.reshape(1,1))
```

```
→ /usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature
  names, but LinearRegression warnings.warn(
array([3.89111601])
```

```
plt.scatter(df['cgpa'],df['package'])
plt.plot(x_train,lr.predict(x_train),color='red')
plt.title('cgpa vs package(in lpa)')
plt.xlabel('cgpa')
plt.ylabel('package(in lpa)')
```

Text(0, 0.5, 'package(in lpa)')



```
m=lr.coef_
b=lr.intercept_
m*8.58+b
```

array([3.89111601])

m*9.5+b

array([4.40443183])

m*100+b

array([54.89908542])

```
from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error
y_pred=lr.predict(x_test)
score=r2_score(y_test,y_pred)
print(f'Accuracy score:{score}')
```

Accuracy score:0.780730147510384

y_pred

array([3.89111601, 3.09324469, 2.38464568, 2.57434935, 1.6537286 ,1.77647803, 2.07219258, 2.93143862, 3.76278706, 2.93701814, 4.09197872, 3.51170867, 2.97049525, 2.40138424, 3.18809652, 3.46707251, 1.94386362, 3.24389172, 2.97607477, 3.41685683, 2.55761079, 3.16577844, 2.85890486, 3.12114229, 3.68467378, 2.8700639 , 3.49497011, 3.34432308, 3.91901361, 1.96060218, 3.65119666, 3.2104146 , 3.74046898, 2.7863711 , 2.78079158, 3.27178932, 3.52844723, 2.61340599, 2.65804215, 2.71383735])

```
MSE=mean_squared_error(y_test,y_pred)
print(f'Mean Squared Error:{MSE}')
```

Mean Squared Error:0.12129235313495527

Multi linear Regression:

```
import numpy as np
import matplotlib.pyplot as
plt
import pandas as pd

df =
pd.read_csv('50_Startups.csv'
) df.head()
```

	R&D Spend	Administration	marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

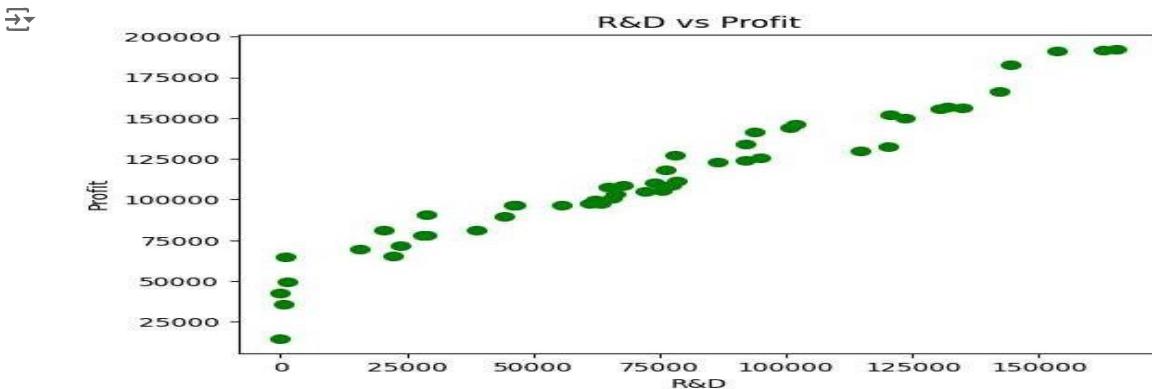
```
df.isnull().sum()
R&D Spend      0
Administration  0
Marketing Spend 0
State          0
Profit         0
dtype: int64

df["State"].unique()

array(['New York', 'California', 'Florida'], dtype=object)

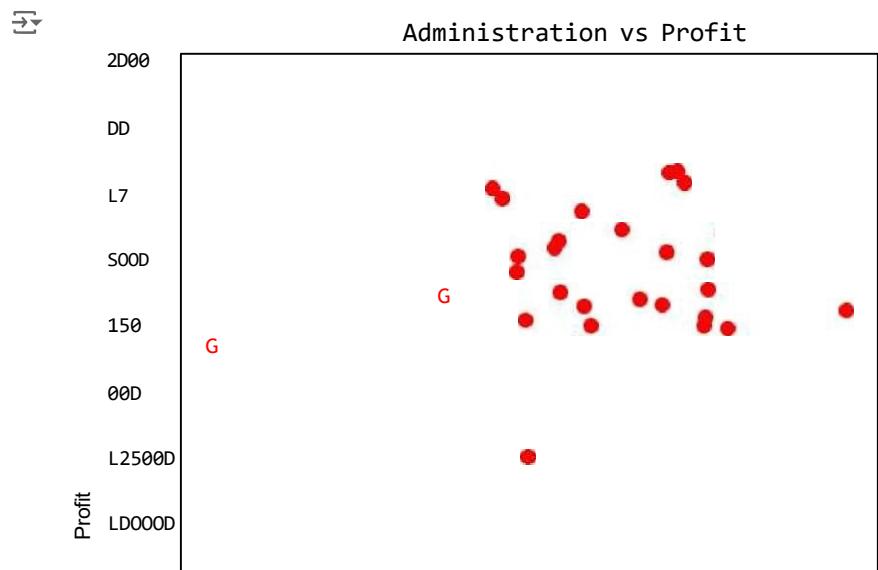
#Plot R&D vs
Profit..... x1 =
df.iloc[:, 0].values
y1 = df.iloc[:, -1].values
plt.scatter(x1,y1,colon='Gree
n',s=50) plt.xlabel('R&D')
plt.ylabel('Profit')
```

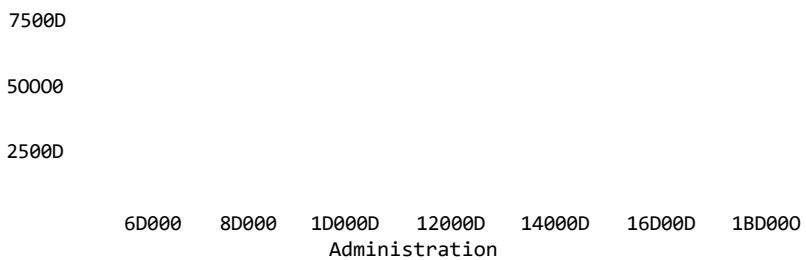
```
plt.title('R&D vs Profit')
plt.show()
```



```
x1 = df.iloc[:, 1].values
y1 = df.iloc[:, -1].values
plt.scatter(x1,y1,color='Red',
s=50)
plt.xlabel('Administration')
plt.ylabel('Profit')
plt.title('Administration vs
Profit')
```

```
plt.show()
```





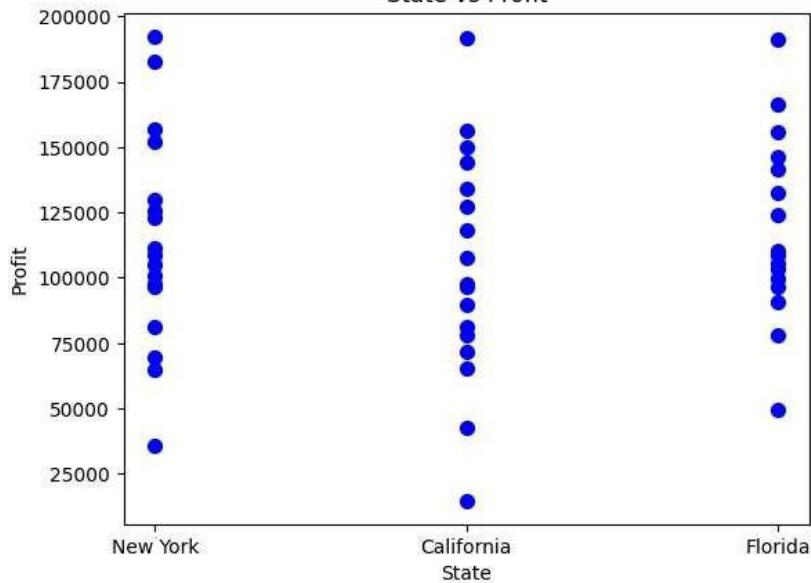
```
x1 = df.iloc[:, 2].values
y1 = df.iloc[:, -1].values
plt.scatter(x1,y1,color='Black',s=50) plt.xlabel('Marketing Spend') plt.ylabel('Profit')
plt.title('Marketing Spend vs Profit') plt.show()
```



```
x1 = df.iloc[:, 3].values
y1 = df.iloc[:, -1].values
plt.scatter(x1,y1,colon='Blue',s=50)
plt.xlabel('State')
plt.ylabel('Profit')
plt.title('State vs Profit') plt.show()
```



State vs Profit



```
df.State.value_coun
```

```
ts()
```

→ State

```
New York    17
California  17
Florida     16
Name: count, dtype: int64
```

```
df['New York'] = np.where(df['State']=='New York', 1, 0)
df['California'] = np.where(df['State']=='California', 1, 0)
df['Florida'] = np.where(df['State']=='Florida', 1, 0)
df.drop(columns=['State'], axis=1, inplace=True)
```

```
df.head()
```

	R&D Spend	AdmInistrat1on	Market1ng Spend	Profit	New York	Ca1iforn1a	Flori da
165349.20	0	136897.80	471784.10	192261.83	1	0	0
162597.70	1	151377.59	443898.53	191792.06	0	1	0
153441.51	2	101145.55	407934.54	191050.39	0	0	1
144372.41	3	118671.85	383199.62	182901.99	1	0	0
142107.34	4	91391.77	366168.42	166187.94	0	0	1

```
y='Profit'
x=df.columns.tolist()
x . remove(y)
```

```

[('R&D Spend', 'Administration', 'Marketing Spend', 'New York', 'California', 'Florida')]
x=df[x].values

y=df[y].values

array([[1.6534920e+05, 1.3689780e+05, 4.7178410e+05, 1.0000000e+00,
       0.0000000e+00, 0.0000000e+00],
      [1.6259770e+05, 1.5137759e+05, 4.4389853e+B5, 0.B000000e+B0,
       1.0000000e+00, 0.0000000e+00],
      [1.5344151e+05, 1.0114555e+05, 4.0793454e+05, 0.0000000e+00,
       0.0000000e+00, 1.0000000e+00],
      [1.4437241e+05, 1.1867185e+05, 3.8319962e+B5, 1.0000000e+B0,
       0.0000000e+00, 0.0000000e+00],
      [1.4210734e+05, 9.1391770e+04, 3.6616842e+05, 0.0000000e+00,
       0.0000000e+00, 1.0000000e+00],
      [1.3187690e+05, 9.9814710e+04, 3.6286136e+05, 1.0000000e+00,
       0.0000000e+00, 0.0000000e+00],
      [1.3461546e+05, 1.4719887e+05, 1.2771682e+05, 8. 8888888e+88,
       1.0000000e+00, 0.0000000e+00],
      [1.3029813e+05, 1.4553006e+05, 3.2387668e+05, 0.0000000e+00,
       0.0000000e+00, 1.0000000e+00],
      [1.2054252e+05, 1.4871895e+05, 3.1161329e+05, 1.0000000e+00,
       0.0000000e+00, 0.0000000e+00],
      [1.2333488e+05, 1.0867917e+05, 3.0498162e+05, 0.0000000e+00,
       1.0000000e+00, 0.0000000e+00],
      [1.0191308e+05, 1.1059411e+B5, 2.2916095e+05, 0.0000000e+00,
       0.0000000e+00, 1.0000000e+00],
      [1.0067196e+05, 9.1790610e+04, 2.4974455e+05, 0.0000000e+00,
       1.0000000e+00, 0.0000000e+00],
      [9.3863750e+04, 1.2732038e+05, 2.4983944e+05, 0.0000000e+00,
       0.0000000e+00, 1.0000000e+00],
      [9.1992390e+04, 1.3549507e+05, 2. 5266493e+05, 1.0000000e+00, 0.0000000e+00],
      [1.1994324e+05, 1.5654742e+05, 2. 5651292e+05,
       0.0000000e+B0, 1.0000000e+00], 0.0000000e+00, 0.0000000e+00,
      [1.1452361e+05, 1.2261684e+05, 2.6177623e+05, 1.0000000e+00,
       0.0000000e+00, 0.0000000e+00],
      [7.8013110e+04, 1.2159755e+05, 2.6434606e+B5, 0.00B0000e+B0,
       1.0000000e+00, 0.0000000e+00],
      [9.4657160e+04, 1.4507758e+05, 2.8257431e+05, 1.0000000e+00,
       0.0000000e+00, 0.0000000e+00], 0.0000000e+00, 1.0000000e+00],
      [8.6419700e+04, 1.5351411e+05, 0.0000000e+00, 1.0000000e+00,
       0.0000000e+00, 0.0000000e+00],
      [7.6253860e+04, 1.1386730e+05, 2.9866447e+05, 0.0000000e+00,
       1.0000000e+00, 0.0000000e+00],
      [7.8389470e+04, 1.5377343e+05, 2.9973729e+05, 1.0000000e+00,
       0.0000000e+00, 0.0000000e+00],
      [7.3994560e+04, 1.2278275e+05, 3.0331926e+05, 0.0000000e+00,
       0.0000000e+00, 1.0000000e+00],
      [6.7532530e+04, 1.0575103e+05, 3.0476873e+05, 0.0000000e+00,
       0.0000000e+00, 1.0000000e+00],
      [7.7044010e+04, 9.9281340e+04, 1.4057481e+05, 1.0000000e+00,
       0.0000000e+00, 0.0000000e+00],
      [6.4664710e+04, 1.3955316e+05, 1.3796262e+05, 0.0000000e+00,
       1.0000000e+00, 0.0000000e+00],
      [7.5328870e+04, 1.4413598e+05, 1.3405007e+05, 0.0000000e+00,
       0.0000000e+00, 1.0000000e+00],
      [7.2107600e+04, 1.2786455e+05, 3.5318381e+05, 1.0000000e+00,
       0.0000000e+00, 0.0000000e+80],
      [6.6851520e+04, 1.8264556e+05, 1.1814820e+B5, 0.0000000e+B0,
       0.0000000e+00, 1.0000000e+00]],

```

y

```
array([192261.83, 191792.06, 191050.39, 182901.99, 166187.94, 156991.12,
       156122.51, 155752.6, 152211.77, 149759.96, 146121.95, 144259.4,
       141585.52, 134307.35, 132602.65, 129917.04, 126992.93, 125370.37,
       124266.9, 122776.86, 118474.03, 111313.02, 110352.25, 108733.99,
       108552.04, 107404.34, 105733.54, 105008.31, 103282.38, 101004.64,
       99937.59, 97483.56, 97427.84, 96778.92, 96712.8, 96479.51,
       90708.19, 89949.14, 81229.06, 81005.76, 78239.91, 77798.83,
       71498.49, 69758.98, 65200.33, 64926.08, 49496.75, 42559.73,
       35673.41, 14681.4 ])
```

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
X_train,X_test,y_train,y_test =
train_test_split(x,y,test_size=0.2,random_state=2)
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
```

```
#Transforming Data
scaler =
MinMaxScaler()
X_train =
scaler.fit_transform(X_train)
X_test =
scaler.fit_transform(X_test)
```

X_train

>Show hidden output

```
lr = LinearRegression()
```

```
lr.fit(X_train,y_train)
```

```
-] * LinearRegression
LinearRegression()
```

```
lr.predict(X_test)
```

```
array([ 74690.00101189, 47320.3206191, 99717.11382077, 155026.3673854, 127299.64727389, 191081.36739932,
64811.64693661, 55816.26398031,
85207.83029741, 109464.85584631])
```

```
# Evaluating model
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
```

```
y_pred =
lr.predict(X_test)
score = r2_score(y_test,
y_pred)
print(f'Accuracy Score : {score}')
```

```
-t r Accuracy Score: 0.9782805755154735
```

Logistic Regression:

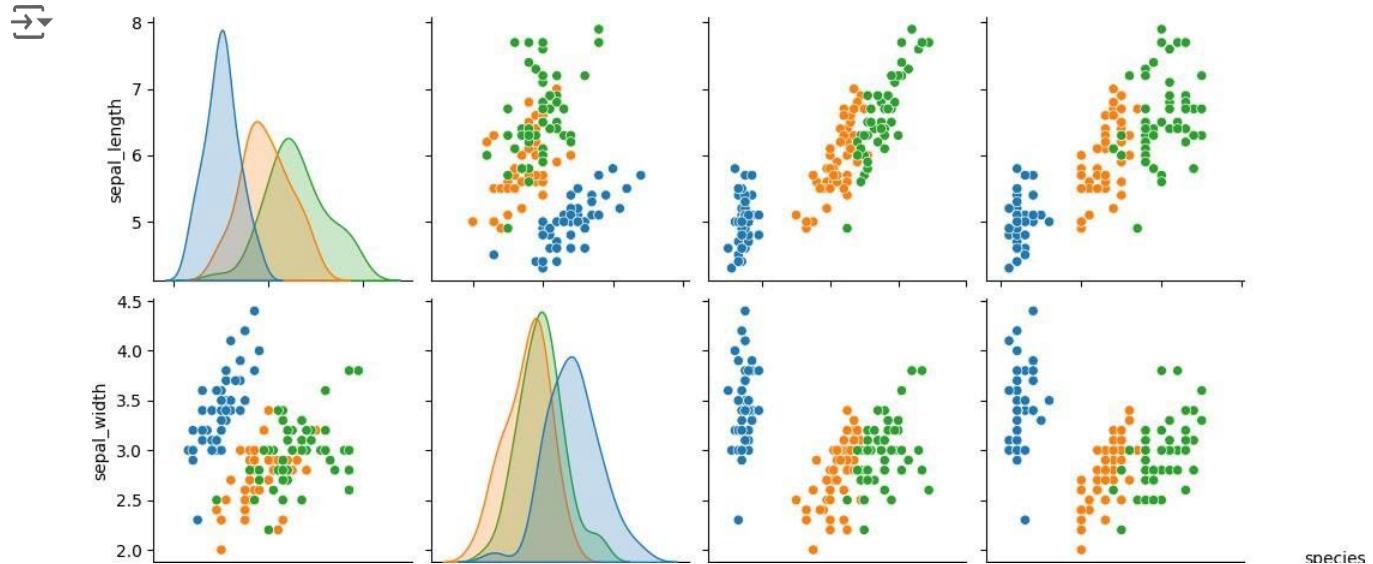
```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
df=sns.load_dataset('iris')
df.head()
```

```
→ sepal_length  sepal_width  petal_length  petal_width  species
```

0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
sns.pairplot(df,hue='species')
plt.show()
```



```
x=df.iloc[:, :-1]
y=df.iloc[:, -1]
```

```
x
```

```
sepal_length  sepal_width  petal_length  petal_width
0           5.1          3.5         1.4        0.2
1           4.9          3.0         1.4        0.2
2           4.7          3.2         1.3        0.2
3           4.6          3.1         1.5        0.2
4           5.0          3.6         1.4        0.2
...
145          6.7          3.0         5.2        2.3
146          6.3          2.5         5.0        1.9
147          6.5          3.0         5.2        2.0
148          6.2          3.4         5.4        2.3
149          5.9          3.0         5.1        1.8
```

150 rows × 4 columns

y

```
setosa
setosa
setosa
setosa
setosa
...
virginica
virginica
virginica
virginica
virginica
Name: species, Length: 150, dtype: object
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=42)

from sklearn.linear_model import LogisticRegression
classifier=LogisticRegression()
classifier.fit(x_train,y_train)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458:
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown
in: <https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
▼
```

```
LogisticRegression
```

```
classifier.predict(x_test)

→ array(['versicolor', 'setosa', 'virginica', 'versicolor',
       'versicolor', 'setosa', 'versicolor', 'virginica',
       'versicolor', 'versicolor',
       'virginica', 'setosa', 'setosa', 'setosa', 'setosa',
       'versicolor', 'virginica', 'versicolor', 'versicolor',
       'virginica', 'setosa',
       'virginica', 'setosa', 'virginica', 'virginica',
       'virginica',
       'virginica', 'virginica', 'setosa', 'setosa', 'setosa',
       'setosa', 'versicolor', 'setosa', 'setosa', 'virginica',
       'versicolor',
       'setosa'], dtype=object)
```

```
from sklearn.metrics import accuracy_score, classification_report
y_pred=classifier.predict(x_test)
score=accuracy_score(y_test,y_pred)
print(score-0.08)
```

```
→ 0.92
```

Experiment - 4

Aim: Implement the naïve Bayesian Classifier for a sample training data set stored as a .CSV file. Compute the accuracy of the classifier, considering a few test data sets.

Naïve Bayesian Classifier:

```
import pandas as pd

df = pd.read_csv('/content/titanic.csv')
df.head()

>Show hidden output

df.drop(['PassengerId', 'Name', 'SibSp', 'Parch', 'Ticket', 'Cabin', 'Embarked'], axis='columns', inplace=True)
df.head()

>Show hidden output

      Pclass   Sex   Age     Fare  Survived
0        3  male  22.0    7.2500       0
1        1 female  38.0   71.2833       1
2        3 female  26.0    7.9250       1
3        1 female  35.0   53.1000       1
4        3  male  35.0    8.0500       0

inputs = df.drop('Survived', axis='columns')
target = df.Survived

inputs.Sex = inputs.Sex.map({'male': 1, 'female': 2})

# dummies = pd.get_dummies(inputs.Sex)      # get dummies method used for mapping
# dummies.head(3)

>Show hidden output

      female   male
0    False   True
1    True  False
2   True  False

# inputs = pd.concat([inputs, dummies], axis='columns')      # concat method used for merging
# inputs.head(3)

>Show hidden output

      Pclass  Sex   Age     Fare  female  male
0        3    1  22.0    7.2500   False  True
1        1    2  38.0   71.2833   True  False
2        3    2  26.0    7.9250   True  False

>Show hidden output

# here we are dropping male column as cause of dummy variable trap theory, One column is enough to represent male vs female
# inputs.drop(['Sex', 'male'], axis='columns', inplace=True)
# inputs.head(3)

inputs.columns[inputs.isna().any()]

>Show hidden output

Index(['Age'], dtype='object')

inputs.Age[:10]
```

Age

```
0 22.0
1 38.0
2 26.0
3 35.0
4 35.0
5 NaN
6 54.0
7 2.0
8 27.0
9 14.0
```

```
inputs.Age = inputs.Age.fillna(inputs.Age.mean())
inputs.head()
```

	Pclass	Sex	Age	Fare
0	3	1	22.0	7.2500
1	1	2	38.0	71.2833
2	3	2	26.0	7.9250
3	1	2	35.0	53.1000
4	3	1	35.0	8.0500

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(inputs, target, test_size=0.3)
```

```
from sklearn.naive_bayes import GaussianNB
model = GaussianNB()
```

```
model.fit(X_train, y_train)
```

```
+ GaussianNB
GaussianNB()
```

```
model.score(X_test, y_test)
```

```
0.7798507462686567
```

```
# calculating the score using cross validation
# from sklearn.model_selection import cross_val_score
# cross_val_score(GaussianNB(), X_train, y_train, cv=5)
```

```
from sklearn.naive_bayes import BernoulliNB
model1 = BernoulliNB()
```

```
model1.fit(X_train, y_train)
```

```
+ BernoulliNB
BernoulliNB()
```

```
model1.score(X_test, y_test)
```

```
0.5895522388059702
```

```
from sklearn.linear_model import LogisticRegression
model2 = LogisticRegression()
```

```
model2.fit(X_train, y_train)
```

```
↳ [+] LogisticRegression  
LogisticRegression()
```

```
model2.score(X_test, y_test)
```

```
↳ 0.7947761194029851
```

Experiment - 5

Aim: Assuming a set of documents that need to be classified, use the Naïve Bayesian Classifier model to perform this task.

Naïve Bayesian Classifier:

```
import pandas as pd

df = pd.read_csv("/content/spam.csv")
df.head()



|   | Category | Message                                           |
|---|----------|---------------------------------------------------|
| 0 | ham      | Go until jurong point, crazy.. Available only ... |
| 1 | ham      | Ok lar... Joking wif u oni...                     |
| 2 | spam     | Free entry in 2 a wkly comp to win FA Cup fina... |
| 3 | ham      | U dun say so early hor... U c already then say... |
| 4 | ham      | Nah I don't think he goes to usf, he lives aro... |



df.groupby('Category').describe()



| Category | Message |        |                                                   | freq |
|----------|---------|--------|---------------------------------------------------|------|
|          | count   | unique | top                                               |      |
| ham      | 4825    | 4516   | Sorry, I'll call later                            | 30   |
| spam     | 747     | 641    | Please call our customer service representativ... | 4    |



df['spam'] = df['Category'].apply(lambda x: 1 if x=='spam' else 0)
df.head()



|   | Category | Message                                           | spam |
|---|----------|---------------------------------------------------|------|
| 0 | ham      | Go until jurong point, crazy.. Available only ... | 0    |
| 1 | ham      | Ok lar... Joking wif u oni...                     | 0    |
| 2 | spam     | Free entry in 2 a wkly comp to win FA Cup fina... | 1    |
| 3 | ham      | U dun say so early hor... U c already then say... | 0    |
| 4 | ham      | Nah I don't think he goes to usf, he lives aro... | 0    |



from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df.Message, df.spam, test_size=0.25)

from sklearn.feature_extraction.text import CountVectorizer
v = CountVectorizer() # to know which word has occurred how many times in dataset as it convert it into vector on bases of frequency col
X_train_count = v.fit_transform(X_train.values)
X_train_count.toarray()[:2]



|                                                           |
|-----------------------------------------------------------|
| array([[0, 0, 0, ..., 0, 0, 0], [0, 0, 0, ..., 0, 0, 0]]) |
|-----------------------------------------------------------|



from sklearn.naive_bayes import MultinomialNB
model = MultinomialNB()
model.fit(X_train_count,y_train)



|                 |
|-----------------|
| • MultinomialNB |
| MultinomialNB() |



emails = [
    'Hey mohan, can we get together to watch football game tomorrow?',
    'Upto 20% discount on parking, exclusive offer just for you. Dont miss this reward!'
]
emails_count = v.transform(emails)
model.predict(emails_count)



|               |
|---------------|
| array([0, 1]) |
|---------------|



X_test_count = v.transform(X_test)
model.score(X_test_count, y_test)



|                    |
|--------------------|
| 0.9892318736539842 |
|--------------------|


```

```
# Sklearn Pipeline

from sklearn.pipeline import Pipeline
clf = Pipeline([
    ('vectorizer', CountVectorizer()),
    ('nb', MultinomialNB())
])
clf.fit(X_train, y_train)
```



```
clf.score(X_test, y_test)
→ 0.9892318736539842
```

Experiment - 6

Aim: Decision tree-based ID3 algorithm.

ID3 algorithm:

```
import pandas as pd

df = pd.read_csv("/content/salaries.csv")
df.head()



|   | company | job                 | degree    | salary_more_then_100k |
|---|---------|---------------------|-----------|-----------------------|
| 0 | google  | sales executive     | bachelors | 0                     |
| 1 | google  | sales executive     | masters   | 0                     |
| 2 | google  | business manager    | bachelors | 1                     |
| 3 | google  | business manager    | masters   | 1                     |
| 4 | apple   | computer programmer | bachelors | 0                     |

inputs = df.drop('salary_more_then_100k', axis='columns')

target = df['salary_more_then_100k']

from sklearn.preprocessing import LabelEncoder

le_company = LabelEncoder()          # LabelEncoder method to give labels as they don't have continuous values just like dummies, map me!
le_job = LabelEncoder()
le_degree = LabelEncoder()

inputs['company_n'] = le_company.fit_transform(inputs['company'])
inputs['job_n'] = le_company.fit_transform(inputs['job'])
inputs['degree_n'] = le_company.fit_transform(inputs['degree'])

inputs



|    | company    | job                 | degree    | company_n | job_n | degree_n |
|----|------------|---------------------|-----------|-----------|-------|----------|
| 0  | google     | sales executive     | bachelors | 2         | 2     | 0        |
| 1  | google     | sales executive     | masters   | 2         | 2     | 1        |
| 2  | google     | business manager    | bachelors | 2         | 0     | 0        |
| 3  | google     | business manager    | masters   | 2         | 0     | 1        |
| 4  | google     | computer programmer | bachelors | 2         | 1     | 0        |
| 5  | google     | computer programmer | masters   | 2         | 1     | 1        |
| 6  | abc pharma | sales executive     | masters   | 0         | 2     | 1        |
| 7  | abc pharma | computer programmer | bachelors | 0         | 1     | 0        |
| 8  | abc pharma | business manager    | bachelors | 0         | 0     | 0        |
| 9  | abc pharma | business manager    | masters   | 0         | 0     | 1        |
| 10 | facebook   | sales executive     | bachelors | 1         | 2     | 0        |
| 11 | facebook   | sales executive     | masters   | 1         | 2     | 1        |
| 12 | facebook   | business manager    | bachelors | 1         | 0     | 0        |
| 13 | facebook   | business manager    | masters   | 1         | 0     | 1        |
| 14 | facebook   | computer programmer | bachelors | 1         | 1     | 0        |
| 15 | facebook   | computer programmer | masters   | 1         | 1     | 1        |

inputs_n = inputs.drop(['company', 'job', 'degree'], axis='columns')

inputs_n
```

	company_n	job_n	degree_n
0	2	2	0
1	2	2	1
2	2	0	0
3	2	0	1
4	2	1	0
5	2	1	1
6	0	2	1
7	0	1	0
8	0	0	0
9	0	0	1
10	1	2	0
11	1	2	1
12	1	0	0
13	1	0	1
14	1	1	0
15	1	1	1

target

	salary_more_then_100k
0	0
1	0
2	1
3	1
4	0
5	1
6	0
7	0
8	0
9	1
10	1
11	1
12	1
13	1
14	1
15	1

```
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier

model = tree.DecisionTreeClassifier()

model.fit(inputs_n, target)

model.score(inputs_n, target)

model.predict([[2, 1, 0]])
```

+ DecisionTreeClassifier
DecisionTreeClassifier()

```

→ /usr/local/lib/python3.10/dist-packages/sklearn/base.py:465: UserWarning: X does not have valid feature names, but DecisionTreeClas:
    warnings.warn(
        array([0])

model.predict([[2, 1, 1]])

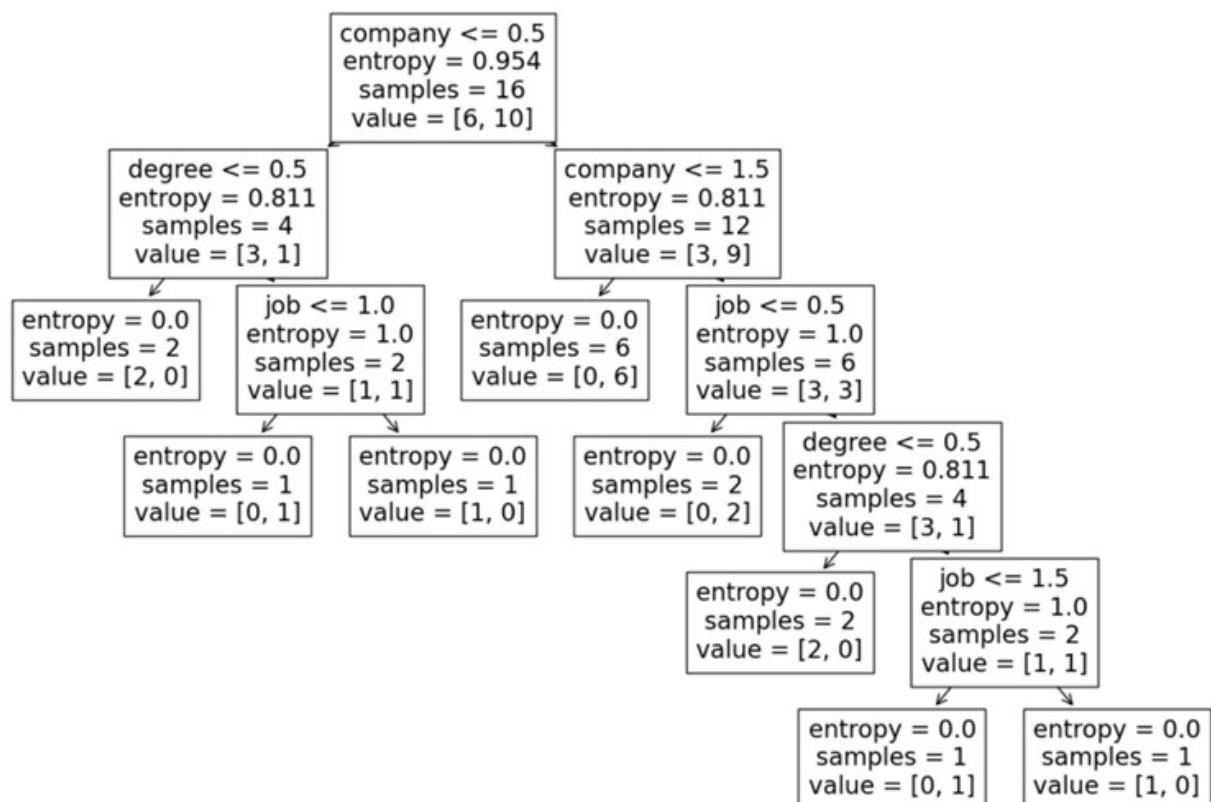
→ /usr/local/lib/python3.10/dist-packages/sklearn/base.py:465: UserWarning: X does not have valid feature names, but DecisionTreeClas:
    warnings.warn(
        array([1])

from sklearn.tree import plot_tree
import matplotlib.pyplot as plt

model = DecisionTreeClassifier(criterion='entropy')
model.fit(inputs_n, target)

plt.figure(figsize=(15, 10))
plot_tree(model, feature_names=df.columns[:-1])
plt.show()

```



Experiment - 7

Aim: Write a program to implement the K- Nearest Neighbour algorithm.

K- Nearest Neighbour:

```
import pandas as pd
from sklearn import datasets
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

cancer_ds = datasets.load_breast_cancer()

x = cancer_ds.data
y = cancer_ds.target

x.shape
→ (569, 30)

y[:5]
→ array([0, 0, 0, 0, 0])

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=0)

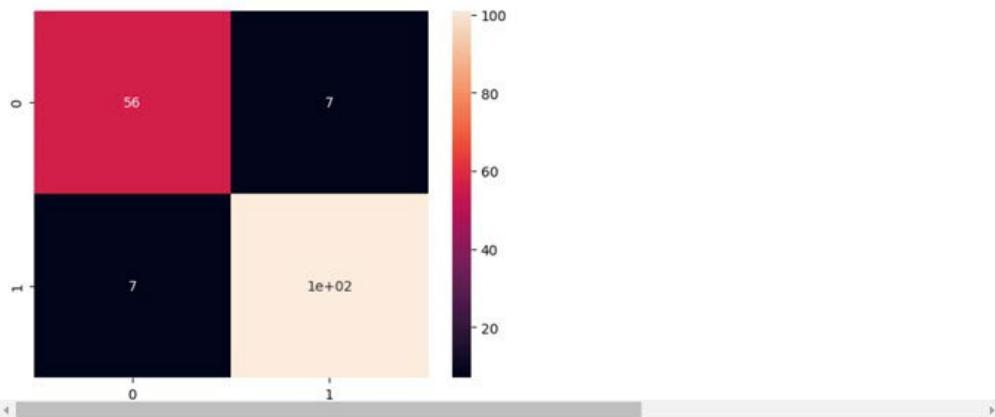
knn_model = KNeighborsClassifier(n_neighbors=3)
knn_model.fit(x_train, y_train)

y_pred = knn_model.predict(x_test)

cm = confusion_matrix(y_test, y_pred)
cls_rpt = classification_report(y_test, y_pred)

print(cm)
sns.heatmap(cm, annot=True)
print(cls_rpt)
```

	precision	recall	f1-score	support
0	0.89	0.89	0.89	63
1	0.94	0.94	0.94	108
accuracy			0.92	171
macro avg	0.91	0.91	0.91	171
weighted avg	0.92	0.92	0.92	171



```
knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(x_train, y_train)

y_pred = knn_model.predict(x_test)
```

```

cm = confusion_matrix(y_test, y_pred)
cls_rpt = classification_report(y_test, y_pred)

print(cm)
print(cls_rpt)

[[ 59  4]
 [ 5 103]]
      precision    recall   f1-score   support
          0       0.92      0.94      0.93      63
          1       0.96      0.95      0.96     108

   accuracy        0.94      0.95      0.94     171
macro avg       0.94      0.95      0.94     171
weighted avg    0.95      0.95      0.95     171

x_train.shape
(398, 30)

knn_model = KNeighborsClassifier(n_neighbors=6)
knn_model.fit(x_train, y_train)

y_pred = knn_model.predict(x_test)

cm = confusion_matrix(y_test, y_pred)
cls_rpt = classification_report(y_test, y_pred)

print(cm)
print(cls_rpt)

[[ 59  4]
 [ 7 101]]
      precision    recall   f1-score   support
          0       0.89      0.94      0.91      63
          1       0.96      0.94      0.95     108

   accuracy        0.93      0.94      0.93     171
macro avg       0.93      0.94      0.93     171
weighted avg    0.94      0.94      0.94     171

knn_model = KNeighborsClassifier(n_neighbors=9)
knn_model.fit(x_train, y_train)

y_pred = knn_model.predict(x_test)

cm = confusion_matrix(y_test, y_pred)
cls_rpt = classification_report(y_test, y_pred)

print(cm)
print(cls_rpt)

[[ 59  4]
 [ 3 105]]
      precision    recall   f1-score   support
          0       0.95      0.94      0.94      63
          1       0.96      0.97      0.97     108

   accuracy        0.96      0.95      0.96     171
macro avg       0.96      0.95      0.96     171
weighted avg    0.96      0.96      0.96     171

knn_model = KNeighborsClassifier(n_neighbors=11)
knn_model.fit(x_train, y_train)

y_pred = knn_model.predict(x_test)

cm = confusion_matrix(y_test, y_pred)
cls_rpt = classification_report(y_test, y_pred)

print(cm)
print(cls_rpt)

```

```

[[ 59  4]
 [ 2 106]]
precision    recall   f1-score   support
          0       0.97      0.94      0.95       63
          1       0.96      0.98      0.97      108
accuracy                           0.96      171
macro avg       0.97      0.96      0.96      171
weighted avg     0.96      0.96      0.96      171

# hyper parameter tuning

from sklearn.model_selection import GridSearchCV

knn_model = KNeighborsClassifier()
params = {'n_neighbors':[3, 7, 9, 11, 13, 17, 19, 21]}

gscv = GridSearchCV(knn_model, param_grid=params, scoring='accuracy')
gscv.fit(x_train, y_train)

GridSearchCV
+ estimator: KNeighborsClassifier
  + KNeighborsClassifier

gscv.best_score_
0.9346518987341772

gscv.best_params_
{'n_neighbors': 11}

model = gscv.best_estimator_

y_pred = model.predict(x_test)
cm = confusion_matrix(y_test, y_pred)
cls_rpt = classification_report(y_test, y_pred)

print(cm)
print(cls_rpt)

[[ 59  4]
 [ 2 106]]
precision    recall   f1-score   support
          0       0.97      0.94      0.95       63
          1       0.96      0.98      0.97      108
accuracy                           0.96      171
macro avg       0.97      0.96      0.96      171
weighted avg     0.96      0.96      0.96      171

```

Experiment - 8

Aim: Apply EM algorithm to cluster a set of data stored in a .CSV file. Use the same data set for clustering using K- Means Algorithm.

Clustering using K- Means Algorithm:

```
# Unsupervised learning
from sklearn import datasets
from sklearn.cluster import KMeans
from matplotlib import pyplot as plt

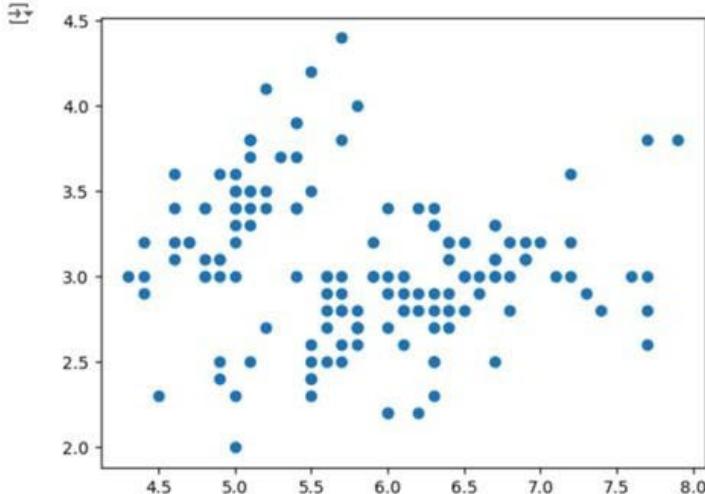
iris_ds = datasets.load_iris()
x = iris_ds.data    # independent variable
x.shape

(150, 4)

# iris_ds

x[:1]      # value of first row
array([[5.1, 3.5, 1.4, 0.2]])

plt.scatter(x[:,0], x[:,1])    # (x[:,0], x[:,1]) = (x[all rows, 1st column], x[all rows, 2nd column])
plt.show()
```



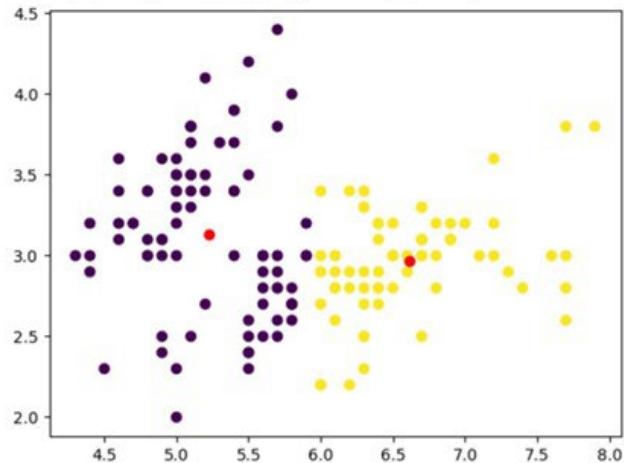
```
cluster_2_model = KMeans(n_clusters=2)
cluster_labels = cluster_2_model.fit_predict(x[:,[0,1]])

cluster_labels[:10], cluster_labels[-10:]

/usr/local/lib/python3.10/dist-packages/scikit-learn/_kmeans.py:1416: FutureWarning: The default value of 'n_init' will change from 10 to 3 in 0.23 (sk-learn v1.1), for consistency with scikit-learn's other clustering estimators.
  super().__check_params_vs_input(X, default_n_init=10)
(array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int32),
 array([1, 1, 0, 1, 1, 1, 1, 1, 1, 0], dtype=int32))

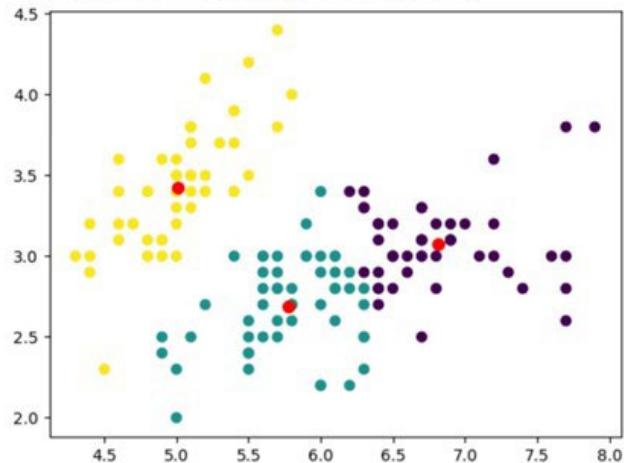
cluster_2_model = KMeans(n_clusters=2)
cluster_labels = cluster_2_model.fit_predict(x[:,[0,1]])
plt.scatter(x[:,0], x[:,1], c=cluster_labels)
centroids = cluster_2_model.cluster_centers_
plt.scatter(centroids[:,0], centroids[:,1], c='red')
plt.show()
```

```
→ /usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:1416: FutureWarning: The default value of 'n_init' will change from 10 to 20 in 0.23. To suppress this warning, use n_init=10 or n_init='auto'. See https://scikit-learn.org/stable/whats_new/0.23.html#deprecations
  super().__check_params_vs_input(X, default_n_init=10)
```



```
cluster_3_model = KMeans(n_clusters=3)
cluster_labels = cluster_3_model.fit_predict(x[:,[0,1]])
plt.scatter(x[:,0], x[:,1], c=cluster_labels)
centroids = cluster_3_model.cluster_centers_
plt.scatter(centroids[:,0], centroids[:,1], s=50, color='red')
plt.show()
```

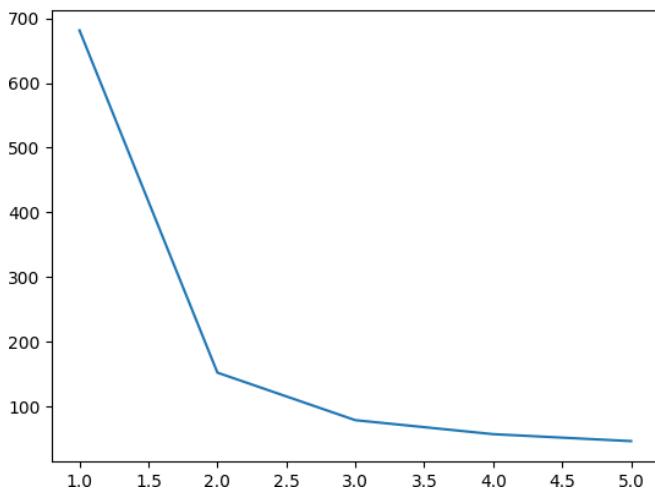
```
→ /usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:1416: FutureWarning: The default value of 'n_init' will change from 10 to 20 in 0.23. To suppress this warning, use n_init=10 or n_init='auto'. See https://scikit-learn.org/stable/whats_new/0.23.html#deprecations
  super().__check_params_vs_input(X, default_n_init=10)
```



```
k_values = [1, 2, 3, 4, 5]
wcss_values = []

for k_value in k_values:
    model = KMeans(n_clusters=k_value)
    model.fit(x)
    wcss_value = model.inertia_
    wcss_values.append(wcss_value)

plt.plot(k_values, wcss_values)
plt.show()
```



```
from sklearn.metrics import silhouette_score

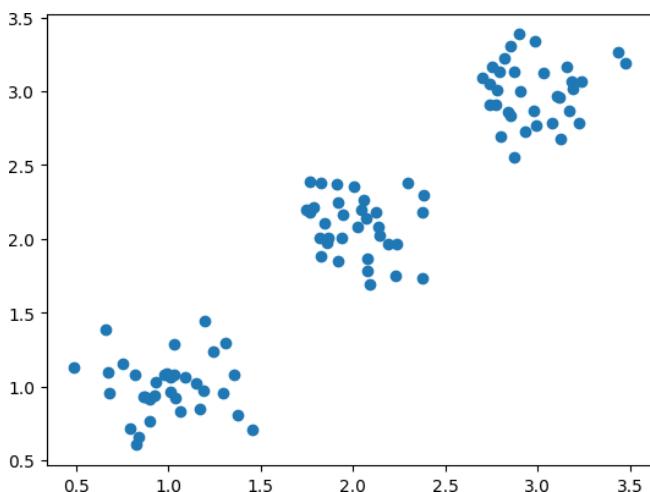
k_values = [2,3,4,5]
wcss_values = []

for k_value in k_values:
    model = KMeans(n_clusters=k_value)
    cluster_labels = model.fit_predict(x[:,[0, 1]])
    model_score = silhouette_score(x[:,[0,1]], cluster_labels)

    print(f"Cluster score: {model_score} for {k_value}")

Cluster score: 0.4629549773635977 for 2
Cluster score: 0.4450525692083638 for 3
Cluster score: 0.4206583419478651 for 4
Cluster score: 0.4050931000679723 for 5

x_data = datasets.make_blobs(n_samples = 100, centers=[[1, 1], [2, 2], [3, 3]], cluster_std = 0.2, random_state = 0)
# x_data[0].shape
plt.scatter(x_data[0][:,0], x_data[0][:,1])
plt.show()
```



Experiment - 9

Aim: Write a program to construct a Bayesian network considering medical data. Use this model to demonstrate the diagnosis of heart patients using standard Heart Disease Data Set.

Naïve Bayes:

```
import numpy as np
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
import matplotlib.pyplot as plt
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import AdaBoostClassifier
from sklearn.preprocessing import StandardScaler
import warnings

# Suppress warnings
warnings.filterwarnings('ignore')

patient_data = pd.read_csv('/content/heart_statlog_cleveland_hungary_final.csv')

patient_data.shape
→ (1190, 12)

patient_data.info()

→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1190 entries, 0 to 1189
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   age              1190 non-null    int64  
 1   sex              1190 non-null    int64  
 2   chest pain type 1190 non-null    int64  
 3   resting bp s     1190 non-null    int64  
 4   cholesterol      1190 non-null    int64  
 5   fasting blood sugar 1190 non-null    int64  
 6   resting ecg       1190 non-null    int64  
 7   max heart rate   1190 non-null    int64  
 8   exercise angina   1190 non-null    int64  
 9   oldpeak           1190 non-null    float64 
 10  ST slope          1190 non-null    int64  
 11  target            1190 non-null    int64  
dtypes: float64(1), int64(11)
memory usage: 111.7 KB
```

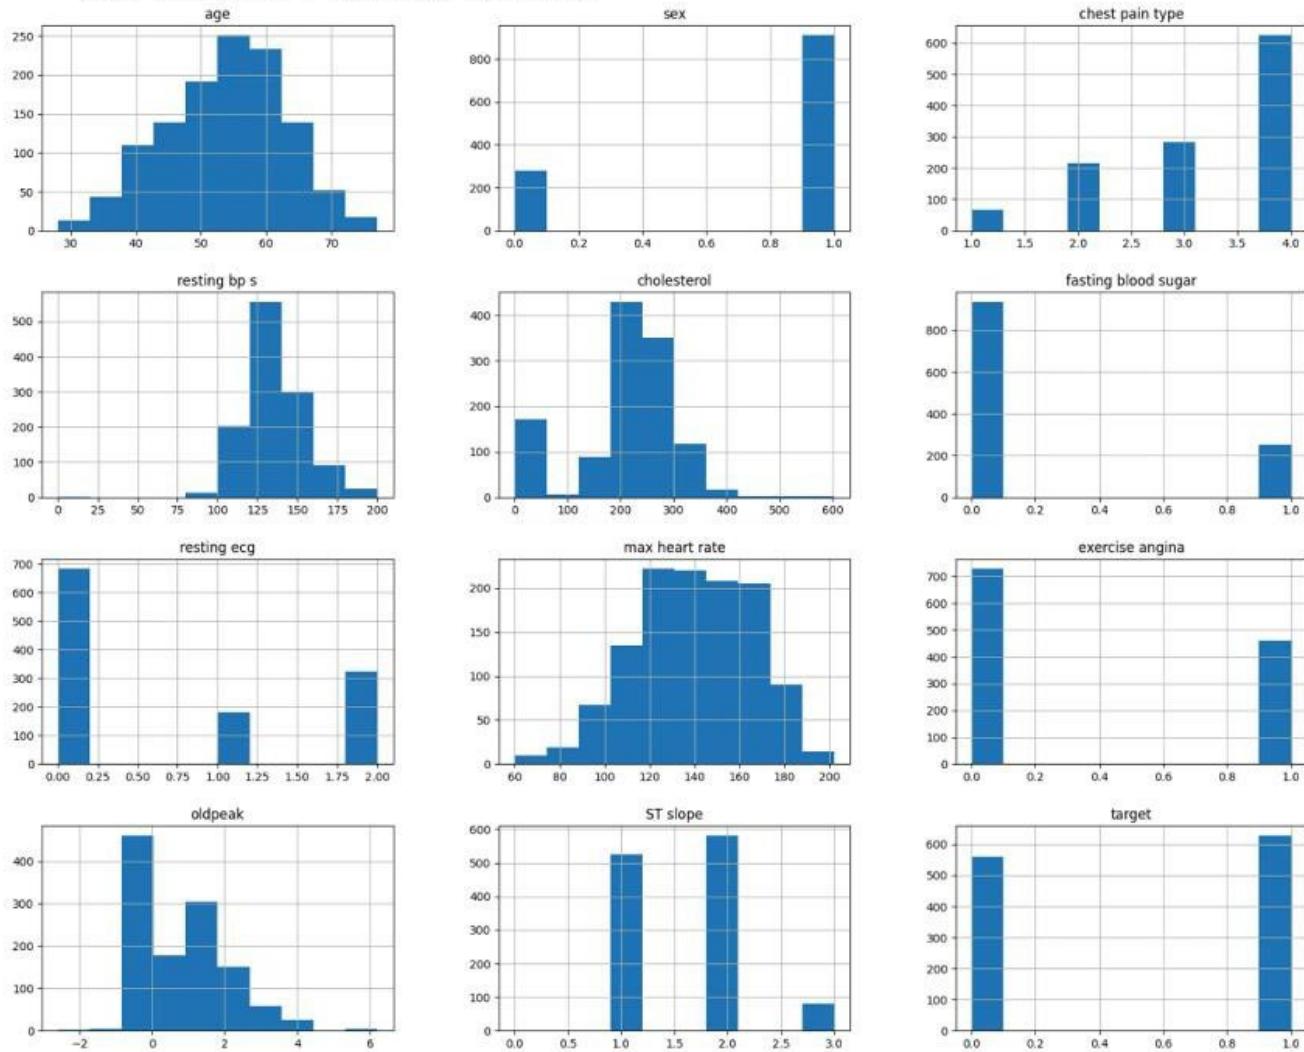
patient_data.head()

	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar	resting ecg	max heart rate	exercise angina	oldpeak	ST slope	target
0	40	1	2	140	289	0	0	172	0	0.0	1	0
1	49	0	3	160	180	0	0	156	0	1.0	2	1
2	37	1	2	130	283	0	1	98	0	0.0	1	0
3	48	0	4	138	214	0	0	108	1	1.5	2	1

```
patient_data.describe()
```

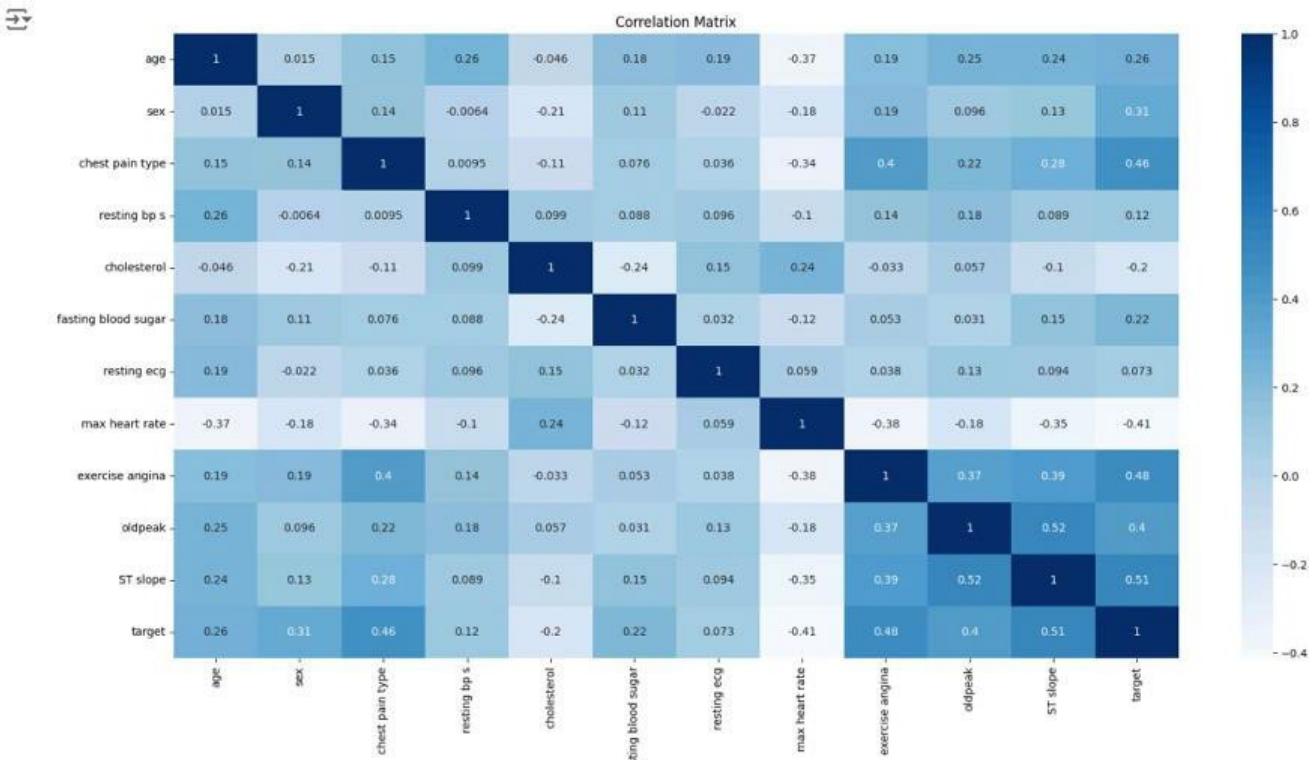
	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar	resting ecg	max heart rate	exercise angina	oldpeak
count	1190.000000	1190.000000	1190.000000	1190.000000	1190.000000	1190.000000	1190.000000	1190.000000	1190.000000	1190.000000
mean	53.720168	0.763866	3.232773	132.153782	210.363866	0.213445	0.698319	139.732773	0.387395	0.922773
std	9.358203	0.424884	0.935480	18.368823	101.420489	0.409912	0.870359	25.517636	0.487360	1.086337
min	28.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	60.000000	0.000000	-2.600000
25%	47.000000	1.000000	3.000000	120.000000	188.000000	0.000000	0.000000	121.000000	0.000000	0.000000
50%	54.000000	1.000000	4.000000	130.000000	229.000000	0.000000	0.000000	140.500000	0.000000	0.600000
75%	60.000000	1.000000	4.000000	140.000000	269.750000	0.000000	2.000000	160.000000	1.000000	1.600000

```
patient_data.hist(figsize=(20, 16))
array([[<Axes: title={'center': 'age'}>, <Axes: title={'center': 'sex'}>,
       <Axes: title={'center': 'chest pain type'}>],
      [<Axes: title={'center': 'resting bp s'}>,
       <Axes: title={'center': 'cholesterol'}>,
       <Axes: title={'center': 'fasting blood sugar'}>],
      [<Axes: title={'center': 'resting ecg'}>,
       <Axes: title={'center': 'max heart rate'}>,
       <Axes: title={'center': 'exercise angina'}>,
       [<Axes: title={'center': 'oldpeak'}>,
        <Axes: title={'center': 'ST slope'}>,
        <Axes: title={'center': 'target'}>]], dtype=object)
```



```
# Heatmap
corr_matrix = patient_data.corr()

plt.figure(figsize=(20, 10))
sns.heatmap(corr_matrix, annot=True, cmap="Blues")
plt.title("Correlation Matrix")
plt.show()
```



```
# Check for null values
patient_data.isnull().sum()
```

	0
age	0
sex	0
chest pain type	0
resting bp s	0
cholesterol	0
fasting blood sugar	0
resting ecg	0
max heart rate	0
exercise angina	0
oldpeak	0
ST slope	0
target	0

```
patient_data.duplicated().sum()
```

```
272
```

```
patient_data["target"].value_counts()
```

```
→ count
target
 1    629
 0    561

# Set target
X_normal = patient_data.drop("target", axis=1)
Y = patient_data["target"]

# Scale data
Scaler = StandardScaler()
X = Scaler.fit_transform(X_normal)
X = pd.DataFrame(X, columns=X_normal.columns)
X.head()

→ age      sex   chest pain type  resting bp s cholesterol   fasting blood sugar  resting ecg  max heart rate  exercise angina  oldpeak  ST slope
 0 -1.466728  0.555995 -1.318351  0.427328  0.775674 -0.520929 -0.802672  1.265039 -0.795219 -0.849792 -1.023217
 1 -0.504600 -1.798576 -0.248932  1.516587 -0.299512 -0.520929 -0.802672  0.637758 -0.795219  0.071119  0.615583
 2 -1.787437  0.555995 -1.318351 -0.117301  0.716489 -0.520929  0.346762 -1.636136 -0.795219 -0.849792 -1.023217
 3 -0.611503 -1.798576  0.820487  0.318402  0.035867 -0.520929 -0.802672 -1.244085  1.257515  0.531575  0.615583
```

Next steps: [Generate code with X](#) [View recommended plots](#) [New interactive sheet](#)

```
Y.value_counts()

→ count
target
 1    629
 0    561

X_train, X_test, Y_train, Y_test = train_test_split(
    X, Y, test_size=0.2, random_state=42
)

# Train and Evaluate model using Naive Bayes model
NB = GaussianNB()

NB.fit(X_train, Y_train)

NB_predict = NB.predict(X_test)

NB_acc_score = accuracy_score(Y_test, NB_predict)

print("Accuracy of Naive Bayes model:", "{:.2f}%".format(NB_acc_score * 100))

print("\nClassification Report:")
nb_cr = classification_report(Y_test, NB_predict)
print(nb_cr)
```

```
NB_cm = confusion_matrix(Y_test, NB_predict)

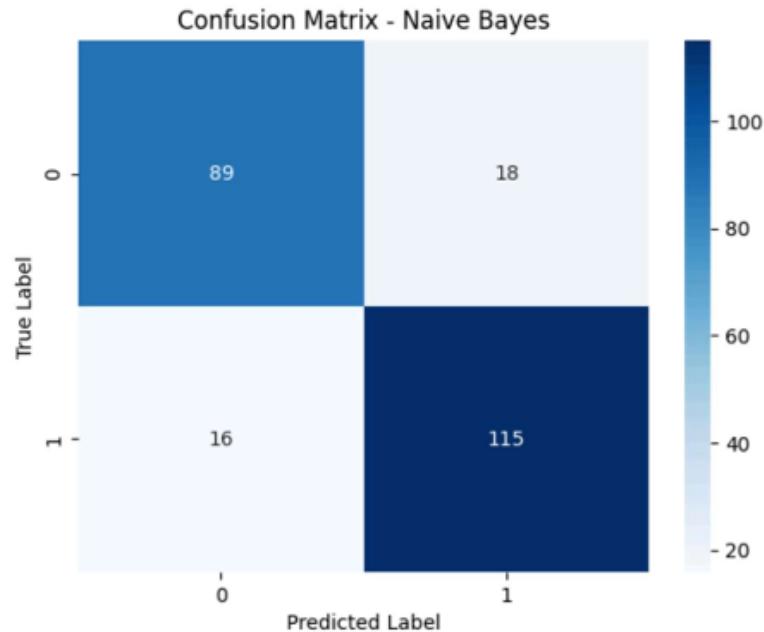
sns.heatmap(NB_cm, annot=True, fmt="d", cmap="Blues")

plt.title("Confusion Matrix - Naive Bayes")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")

plt.show()
```

→ Accuracy of Naive Bayes model: 85.71%

	precision	recall	f1-score	support
0	0.85	0.83	0.84	107
1	0.86	0.88	0.87	131
accuracy			0.86	238
macro avg	0.86	0.85	0.86	238
weighted avg	0.86	0.86	0.86	238



Experiment – 10

Aim: Compare the various Supervised learning algorithm by using appropriate dataset.

Comparison of various Supervised Learning Algorithms:

```
import numpy as np
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
import matplotlib.pyplot as plt
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import AdaBoostClassifier
from sklearn.preprocessing import StandardScaler

import warnings

# Suppress warnings
warnings.filterwarnings("ignore")

patient_data = pd.read_csv("/content/heart_statlog_cleveland_hungary_final.csv")

patient_data.shape
→ (1190, 12)

patient_data.info()

→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1190 entries, 0 to 1189
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   age              1190 non-null    int64  
 1   sex              1190 non-null    int64  
 2   chest pain type 1190 non-null    int64  
 3   resting bp s     1190 non-null    int64  
 4   cholesterol      1190 non-null    int64  
 5   fasting blood sugar 1190 non-null    int64  
 6   resting ecg       1190 non-null    int64  
 7   max heart rate   1190 non-null    int64  
 8   exercise angina  1190 non-null    int64  
 9   oldpeak          1190 non-null    float64 
 10  ST slope         1190 non-null    int64  
 11  target            1190 non-null    int64  
dtypes: float64(1), int64(11)
memory usage: 111.7 KB
```

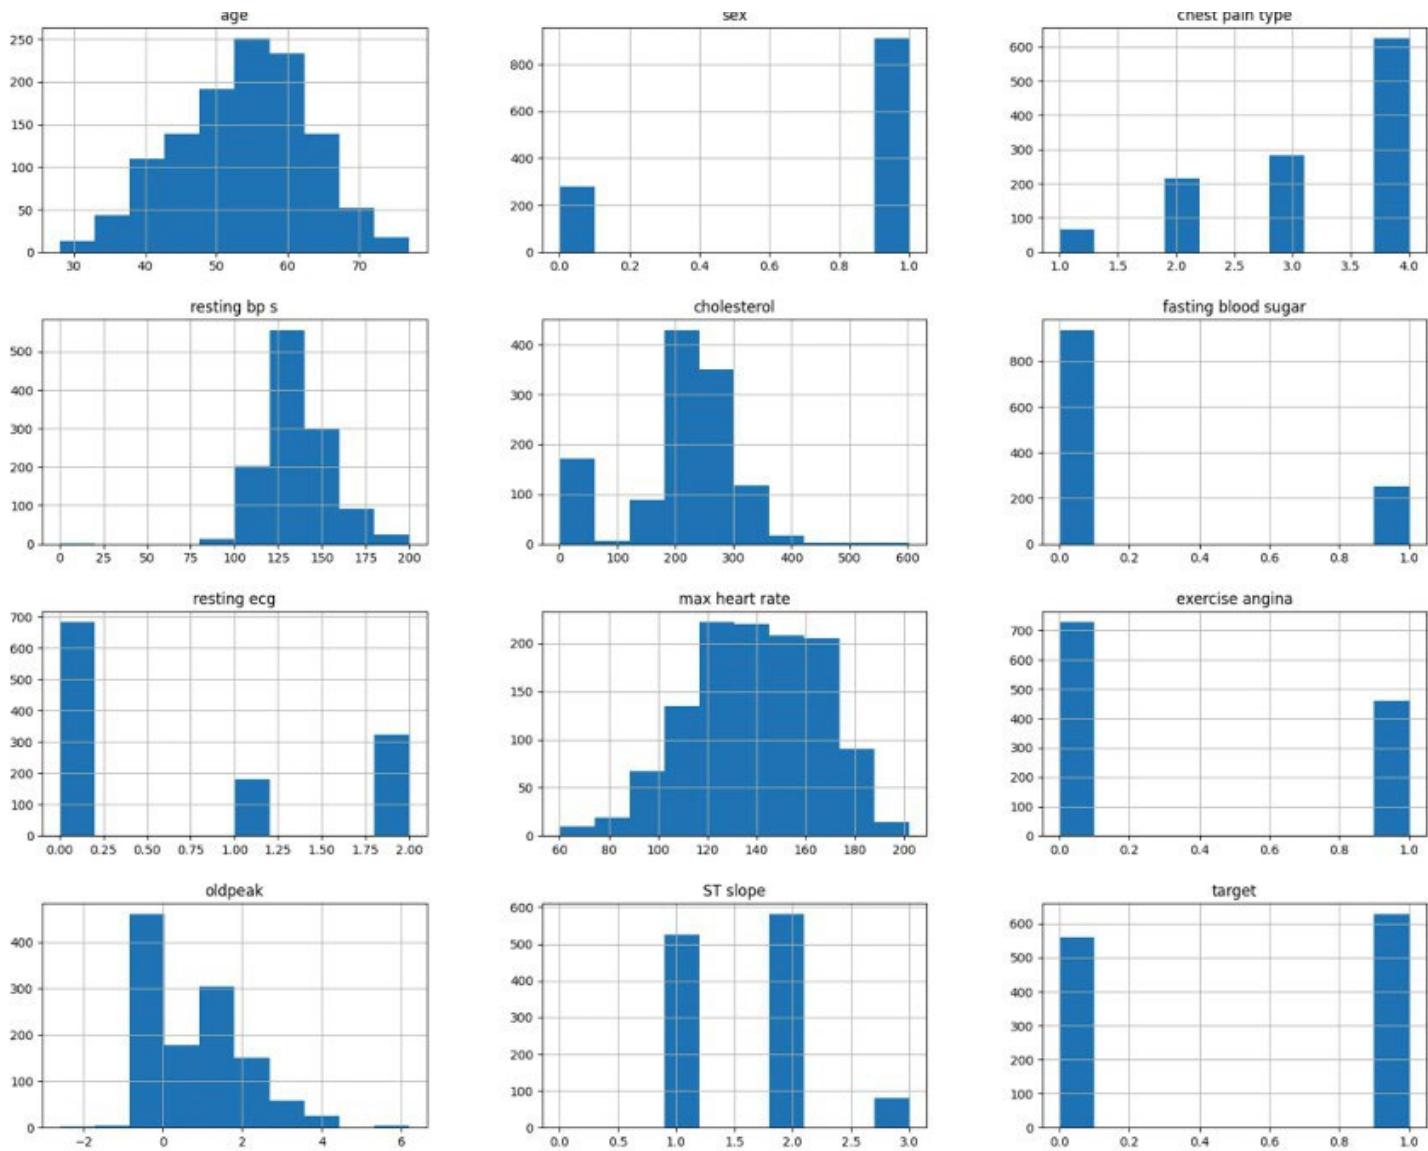
patient_data.head()

	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar	resting ecg	max heart rate	exercise angina	oldpeak	ST slope	target
0	40	1	2	140	289	0	0	172	0	0.0	1	0
1	49	0	3	160	180	0	0	156	0	1.0	2	1
2	37	1	2	130	283	0	1	98	0	0.0	1	0
3	48	0	4	138	214	0	0	108	1	1.5	2	1

patient_data.describe().T

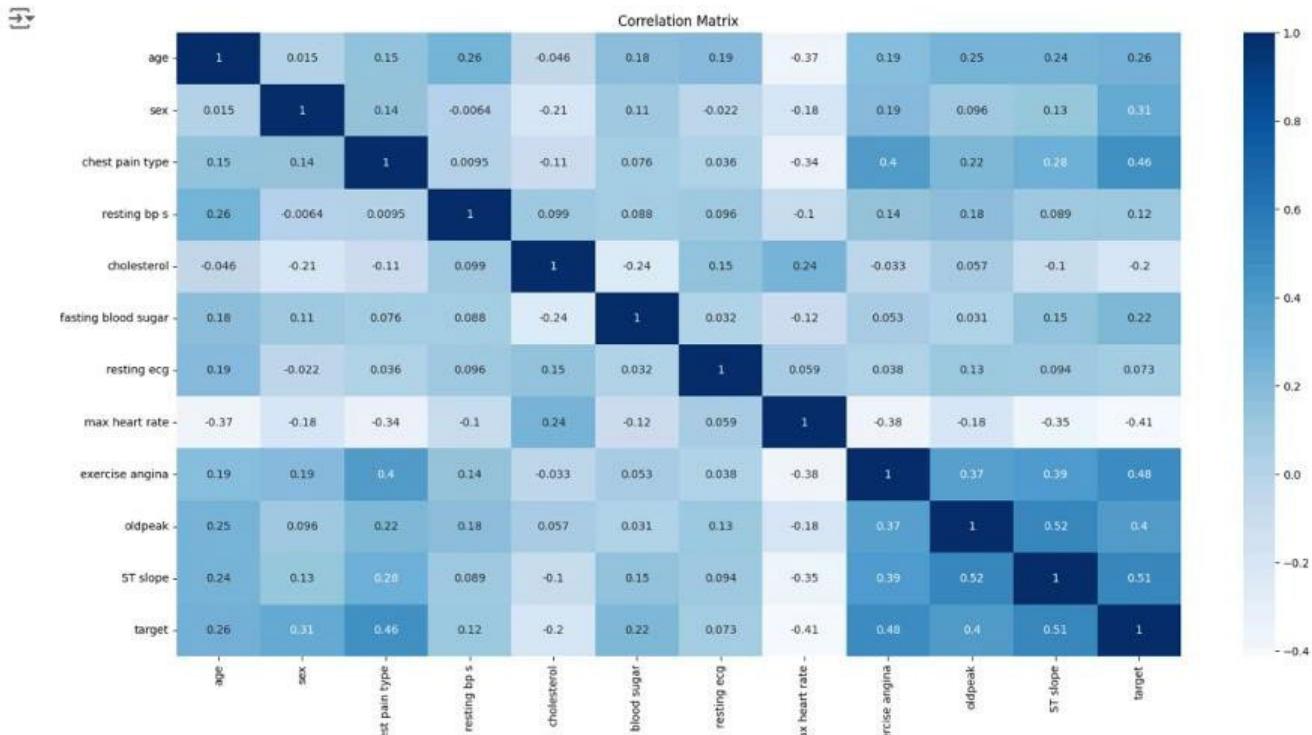
	count	mean	std	min	25%	50%	75%	max	
age	1190.0	53.720168	9.358203	28.0	47.0	54.0	60.00	77.0	
sex	1190.0	0.763866	0.424884	0.0	1.0	1.0	1.00	1.0	
chest pain type	1190.0	3.232773	0.935480	1.0	3.0	4.0	4.00	4.0	
resting bp s	1190.0	132.153782	18.368823	0.0	120.0	130.0	140.00	200.0	
cholesterol	1190.0	210.363866	101.420489	0.0	188.0	229.0	269.75	603.0	
fasting blood sugar	1190.0	0.213445	0.409912	0.0	0.0	0.0	0.00	1.0	
resting ecg	1190.0	0.698319	0.870359	0.0	0.0	0.0	2.00	2.0	
max heart rate	1190.0	139.732773	25.517636	60.0	121.0	140.5	160.00	202.0	
exercise angina	1190.0	0.387395	0.487360	0.0	0.0	0.0	1.00	1.0	
oldpeak	1190.0	0.922773	1.086337	-2.6	0.0	0.6	1.60	6.2	
ST slope	1190.0	1.624370	0.610459	0.0	1.0	2.0	2.00	3.0	
target	1190.0	0.528571	0.499393	0.0	0.0	1.0	1.00	1.0	

patient_data.hist(figsize=(20, 16))



```
corr_matrix = patient_data.corr()

plt.figure(figsize=(20, 10))
sns.heatmap(corr_matrix, annot=True, cmap="Blues")
plt.title("Correlation Matrix")
plt.show()
```



```
patient_data.isnull().sum()
```

	0
age	0
sex	0
chest pain type	0
resting bp s	0
cholesterol	0
fasting blood sugar	0
resting ecg	0
max heart rate	0
exercise angina	0
oldpeak	0
ST slope	0
target	0

```
patient_data.duplicated().sum()
```

```
272
```

```
patient_data["target"].value_counts()
```

count

target

1	629
0	561

```
X_normal = patient_data.drop("target", axis=1)
Y = patient_data["target"]
```

```
Scaler = StandardScaler()
X = Scaler.fit_transform(X_normal)
X = pd.DataFrame(X, columns=X_normal.columns)
X.head()
```

	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar	resting ecg	max heart rate	exercise angina	oldpeak	ST slope
0	-1.466728	0.555995	-1.318351	0.427328	0.775674	-0.520929	-0.802672	1.265039	-0.795219	-0.849792	-1.023217
1	-0.504600	-1.798576	-0.248932	1.516587	-0.299512	-0.520929	-0.802672	0.637758	-0.795219	0.071119	0.615583
2	-1.787437	0.555995	-1.318351	-0.117301	0.716489	-0.520929	0.346762	-1.636136	-0.795219	-0.849792	-1.023217
3	-0.611503	-1.798576	0.820487	0.318402	0.035867	-0.520929	-0.802672	-1.244085	1.257515	0.531575	0.615583

Next steps: [Generate code with X](#) [View recommended plots](#) [New interactive sheet](#)

```
Y.value_counts()
```

count

target

1	629
0	561

```
# 80% of the data will be used for training
# 20% of the data will be used for testing
```

```
X_train, X_test, Y_train, Y_test = train_test_split(
    X, Y, test_size=0.2, random_state=42
)
```

Model training and Evalution using different Supervised Learning Algorithms

1. Logistic Regression

```
# Define the parameter grid for Logistic Regression
logreg_param_grid = {
    "penalty": ["l1", "l2"],
    "C": np.logspace(-3, 3, 7),
    "solver": ["liblinear"],
}
```

```
# Create a Logistic Regression model
logreg_model = LogisticRegression(random_state=42)

# Perform grid search with cross-validation
logreg_grid_search = GridSearchCV(
    logreg_model, logreg_param_grid, cv=5, scoring="accuracy"
)
logreg_grid_search.fit(X_train, Y_train)

# Get the best parameters
best_logreg_params = logreg_grid_search.best_params_
print("Best Hyperparameters for Logistic Regression:", best_logreg_params)

# Train a Logistic Regression model with the best parameters
best_logreg_model = LogisticRegression(random_state=42, **best_logreg_params)
best_logreg_model.fit(X_train, Y_train)

# Make predictions on the test set
logreg_predict = best_logreg_model.predict(X_test)

# Calculate accuracy on the test set
best_logreg_acc = accuracy_score(Y_test, logreg_predict)
print(
    "Best Accuracy of Logistic Regression model:",
    "{:.2f}%".format(best_logreg_acc * 100),
)

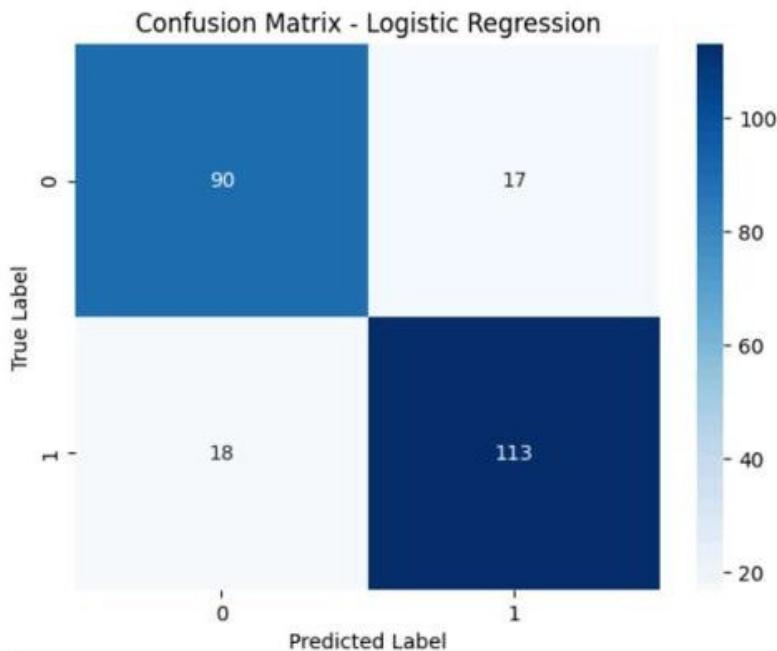
# Display classification report
print("\nClassification Report - Logistic Regression:")
lr_cr = classification_report(Y_test, logreg_predict)
print(lr_cr)

# Display confusion matrix
logreg_cm = confusion_matrix(Y_test, logreg_predict)

# Plot the confusion matrix
sns.heatmap(logreg_cm, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix - Logistic Regression")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

Best Hyperparameters for Logistic Regression: {'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'}
Best Accuracy of Logistic Regression model: 85.29%

Classification Report - Logistic Regression:				
	precision	recall	f1-score	support
0	0.83	0.84	0.84	107
1	0.87	0.86	0.87	131
accuracy			0.85	238
macro avg	0.85	0.85	0.85	238
weighted avg	0.85	0.85	0.85	238



2. Decision Tree

```
# Define the parameter grid
param_grid = {
    "criterion": ["gini", "entropy"],
    "max_depth": np.arange(1, 21),
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 4],
    "max_features": ["sqrt", "log2", None],
}

# Create a Decision Tree model
DT = DecisionTreeClassifier(random_state=0)

# Perform grid search with cross-validation
grid_search = GridSearchCV(DT, param_grid, cv=5, scoring="accuracy")
grid_search.fit(X_train, Y_train)

best_params = grid_search.best_params_
print("Best Hyperparameters:", best_params)

# Train a Decision Tree with the best parameters
best_DT = DecisionTreeClassifier(random_state=0, **best_params)
best_DT.fit(X_train, Y_train)

# Make predictions on the test set
DT_predict = best_DT.predict(X_test)

# Calculate accuracy on the test set
max_dt_acc = accuracy_score(Y_test, DT_predict)
print(
    "Accuracy of Decision Tree with Best Parameters:",
    "{:.2f}%".format(max_dt_acc * 100),
)
```

```

print("\nClassification Report:")
dt_cr = classification_report(Y_test, DT_predict)
print(dt_cr)

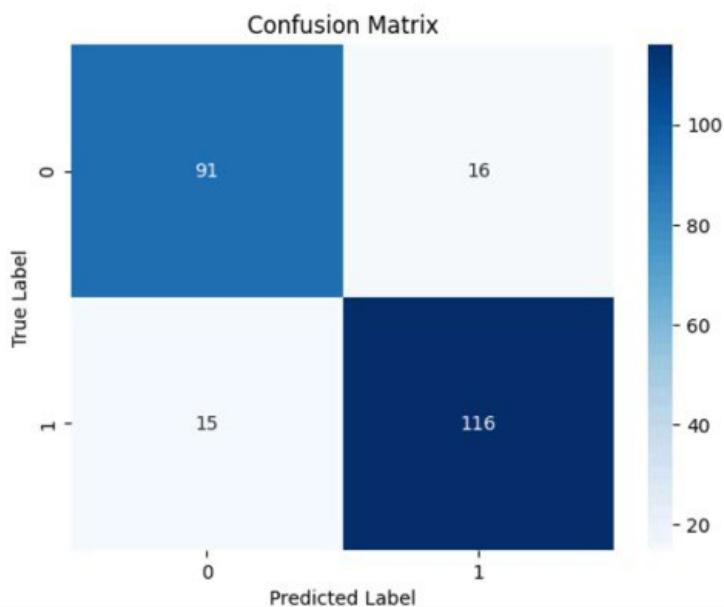
# Display confusion matrix
DT_cm = confusion_matrix(Y_test, DT_predict)

# Plot the confusion matrix
sns.heatmap(DT_cm, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

→ Best Hyperparameters: {'criterion': 'entropy', 'max_depth': 16, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 2}
Accuracy of Decision Tree with Best Parameters: 86.97%

```

Classification Report:				
	precision	recall	f1-score	support
0	0.86	0.85	0.85	107
1	0.88	0.89	0.88	131
accuracy			0.87	238
macro avg	0.87	0.87	0.87	238
weighted avg	0.87	0.87	0.87	238



3. KNN

```

# Define the parameter grid for K-Nearest Neighbors
knn_param_grid = {
    "n_neighbors": np.arange(1, 21),
    "weights": ["uniform", "distance"],
    "algorithm": ["auto", "ball_tree", "kd_tree", "brute"],
}
# Create a K-Nearest Neighbors model
KNN = KNeighborsClassifier()

# Perform grid search with cross-validation
knn_grid_search = GridSearchCV(KNN, knn_param_grid, cv=5, scoring="accuracy")
knn_grid_search.fit(X_train, Y_train)

```

```
best_knn_params = knn_grid_search.best_params_
print("Best Hyperparameters for K-Nearest Neighbors:", best_knn_params)

# Train a K-Nearest Neighbors model with the best parameters
best_KNN_model = KNeighborsClassifier(**best_knn_params)
best_KNN_model.fit(X_train, Y_train)

# Make predictions on the test set
KNN_predict = best_KNN_model.predict(X_test)

# Calculate accuracy on the test set
best_KNN_acc = accuracy_score(Y_test, KNN_predict)
print("Best Accuracy of K-Neighbors Classifier:", "{:.2f}%".format(best_KNN_acc * 100))

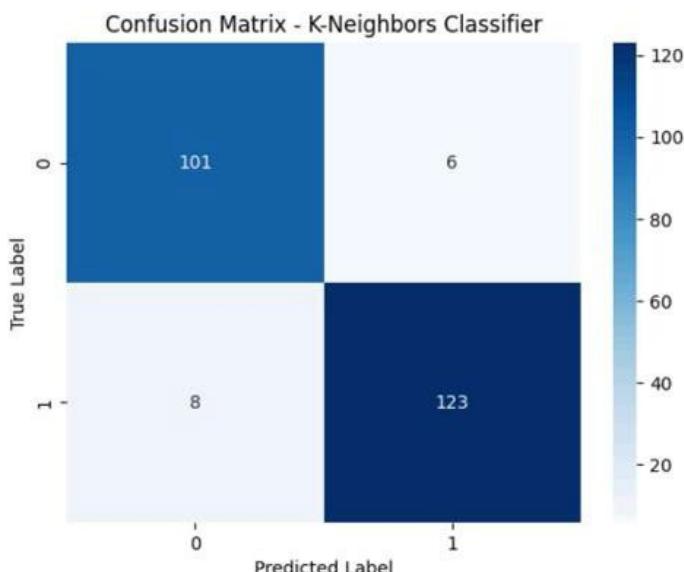
# Display classification report
print("\nClassification Report - K-Neighbors Classifier:")
knn_cr = classification_report(Y_test, KNN_predict)
print(knn_cr)

# Display confusion matrix
KNN_cm = confusion_matrix(Y_test, KNN_predict)

# Plot the confusion matrix
sns.heatmap(KNN_cm, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix - K-Neighbors Classifier")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

→ Best Hyperparameters for K-Nearest Neighbors: {'algorithm': 'auto', 'n_neighbors': 19, 'weights': 'distance'}
Best Accuracy of K-Neighbors Classifier: 94.12%

Classification Report - K-Neighbors Classifier:				
	precision	recall	f1-score	support
0	0.93	0.94	0.94	107
1	0.95	0.94	0.95	131
accuracy			0.94	238
macro avg	0.94	0.94	0.94	238
weighted avg	0.94	0.94	0.94	238



4. SVM

```
# Define the parameter grid for Support Vector Classifier
svm_param_grid = {
    "C": [0.1, 1, 10, 100],
    "gamma": ["scale", "auto", 0.001, 0.01, 0.1, 1, 10],
    "kernel": ["rbf"],
}

# Create a Support Vector Classifier model
SVM = SVC()

# Perform grid search with cross-validation
svm_grid_search = GridSearchCV(SVM, svm_param_grid, cv=5, scoring="accuracy")
svm_grid_search.fit(X_train, Y_train)

# Get the best parameters
best_svm_params = svm_grid_search.best_params_
print("Best Hyperparameters for Support Vector Classifier:", best_svm_params)

# Train a Support Vector Classifier model with the best parameters
best_SVM_model = SVC(**best_svm_params)
best_SVM_model.fit(X_train, Y_train)

# Make predictions on the test set
SVM_predict = best_SVM_model.predict(X_test)

# Calculate accuracy on the test set
best_SVM_acc = accuracy_score(Y_test, SVM_predict)
print(
    "Best Accuracy of Support Vector Classifier:", "{:.2f}%".format(best_SVM_acc * 100)
)

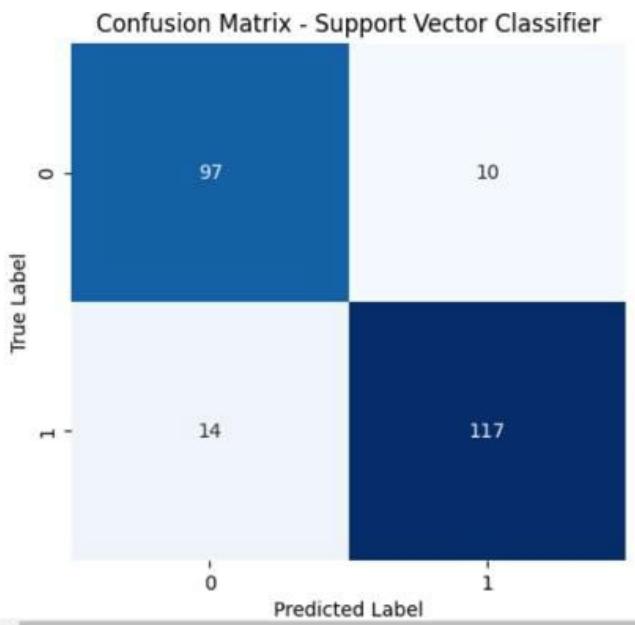
# Display classification report
print("\nClassification Report - Support Vector Classifier:")
svm_cr = classification_report(Y_test, SVM_predict)
print(svm_cr)

# Display confusion matrix
SVM_cm = confusion_matrix(Y_test, SVM_predict)

# Plot the confusion matrix
sns.heatmap(SVM_cm, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix - Support Vector Classifier")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

→ Best Hyperparameters for Support Vector Classifier: {'C': 10, 'gamma': 0.1, 'kernel': 'rbf'}
Best Accuracy of Support Vector Classifier: 89.92%

Classification Report - Support Vector Classifier:
precision    recall   f1-score   support
          0       0.87      0.91      0.89      107
          1       0.92      0.89      0.91      131
   accuracy                           0.90      238
    macro avg       0.90      0.90      0.90      238
 weighted avg       0.90      0.90      0.90      238
```



5. Naive Bayes

```
NB = GaussianNB()
NB.fit(X_train, Y_train)
NB_predict = NB.predict(X_test)

NB_acc_score = accuracy_score(Y_test, NB_predict)

print("Accuracy of Naive Bayes model:", "{:.2f}%".format(NB_acc_score * 100))

print("\nClassification Report:")
nb_cr = classification_report(Y_test, NB_predict)
print(nb_cr)

NB_cm = confusion_matrix(Y_test, NB_predict)

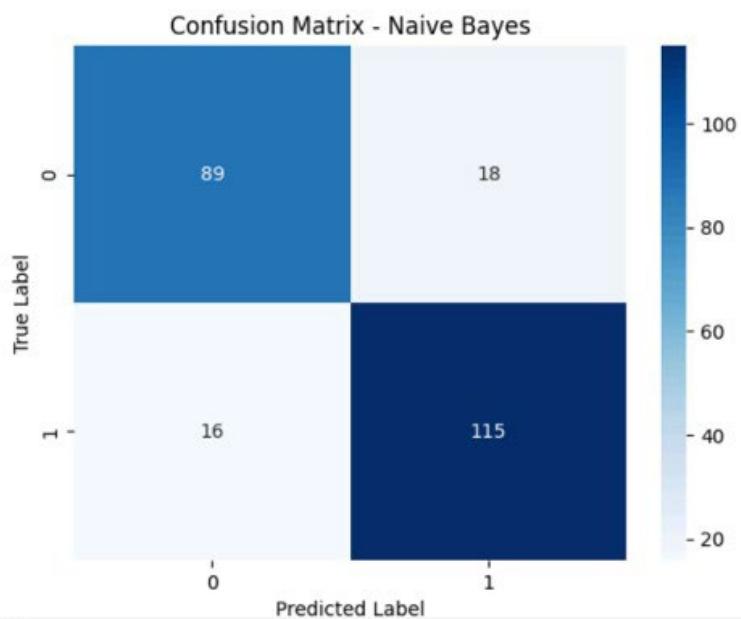
sns.heatmap(NB_cm, annot=True, fmt="d", cmap="Blues")

plt.title("Confusion Matrix - Naive Bayes")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")

plt.show()
```

Accuracy of Naive Bayes model: 85.71%

Classification Report:				
	precision	recall	f1-score	support
0	0.85	0.83	0.84	107
1	0.86	0.88	0.87	131
accuracy			0.86	238
macro avg	0.86	0.85	0.86	238
weighted avg	0.86	0.86	0.86	238



Comparison

1. Comparison Table

```
# Create a dictionary to store the models and their accuracies
models_accuracy = {
    "Logistic Regression": best_logreg_acc,
    "Decision Tree": max_dt_acc,
    "KNN": best_KNN_acc,
    "SVM": best_SVM_acc,
    "Naive Bayes": NB_acc_score,
}

# Find the model with the highest accuracy
best_model = max(models_accuracy, key=models_accuracy.get)
best_accuracy = models_accuracy[best_model]

comparison = pd.DataFrame(
{
    "Model": list(models_accuracy.keys()),
    "Accuracy": ["{:.2f}%".format(acc * 100) for acc in models_accuracy.values()],
}
)

print("Comparison Table:")
print(comparison)

# Print the name and accuracy of the best model
print("\nBest Model:")
print(f"{best_model}: {best_accuracy:.2%}")
```

→ Comparison Table:

Model Accuracy		
0	Logistic Regression	85.29%
1	Decision Tree	86.97%
2	KNN	94.12%
3	SVM	89.92%
4	Naive Bayes	85.71%

Best Model:

KNN: 94.12%

2. Comparison using Bar Plot

```
# Multiply values by 100 and format to 2 decimal places
y_values = [value * 100 for value in models_accuracy.values()]
y_labels = ["{:.2f}%".format(value) for value in y_values]

# Create a color palette with sufficient contrast
colors = sns.color_palette("Set2", len(models_accuracy))

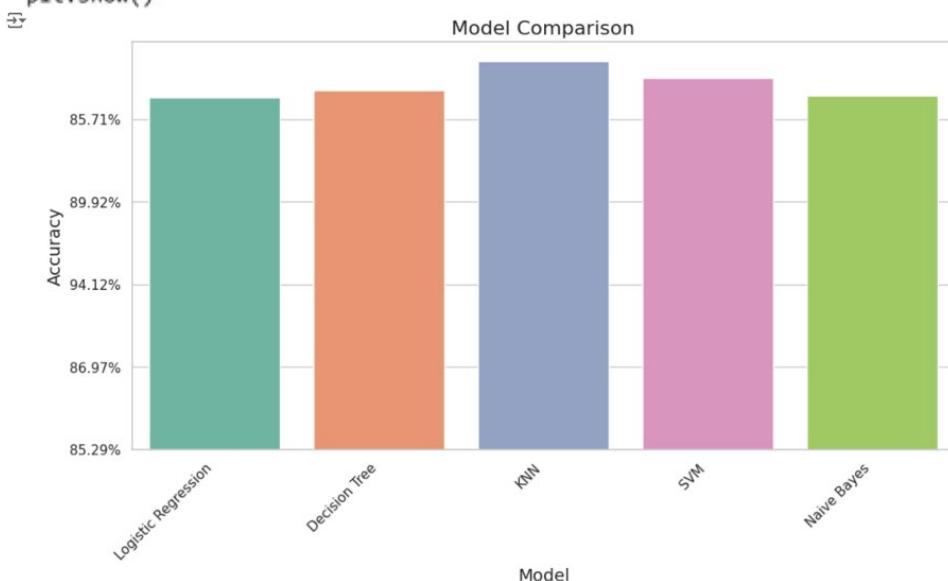
# Create a bar plot
plt.figure(figsize=(12, 6))
sns.set(style="whitegrid")
ax = sns.barplot(x=list(models_accuracy.keys()), y=y_values, palette=colors)

# Set y-axis labels
ax.set_yticklabels(y_labels)

# Rotate x-axis labels for better readability
plt.xticks(rotation=45, ha="right")

plt.title("Model Comparison", fontsize=16)
plt.xlabel("Model", fontsize=14)
plt.ylabel("Accuracy", fontsize=14)

plt.show()
```



3. Comparing Confusion Matrices

```

num_classifiers = 5
num_rows = (num_classifiers - 1) // 4 + 1
num_cols = min(num_classifiers, 4)

fig, axes = plt.subplots(
    nrows=num_rows, ncols=num_cols, figsize=(5 * num_cols, 5 * num_rows)
)
fig.suptitle("Confusion Matrices", fontsize=20)

classifiers = [
    ("Logistic Regression", logreg_cm, best_logreg_acc),
    ("Decision Tree", DT_cm, max_dt_acc),
    ("K-Neighbors", KNN_cm, best_KNN_acc),
    ("SVM", SVM_cm, best_SVM_acc),
    ("Naive Bayes", NB_cm, NB_acc_score),
]
]

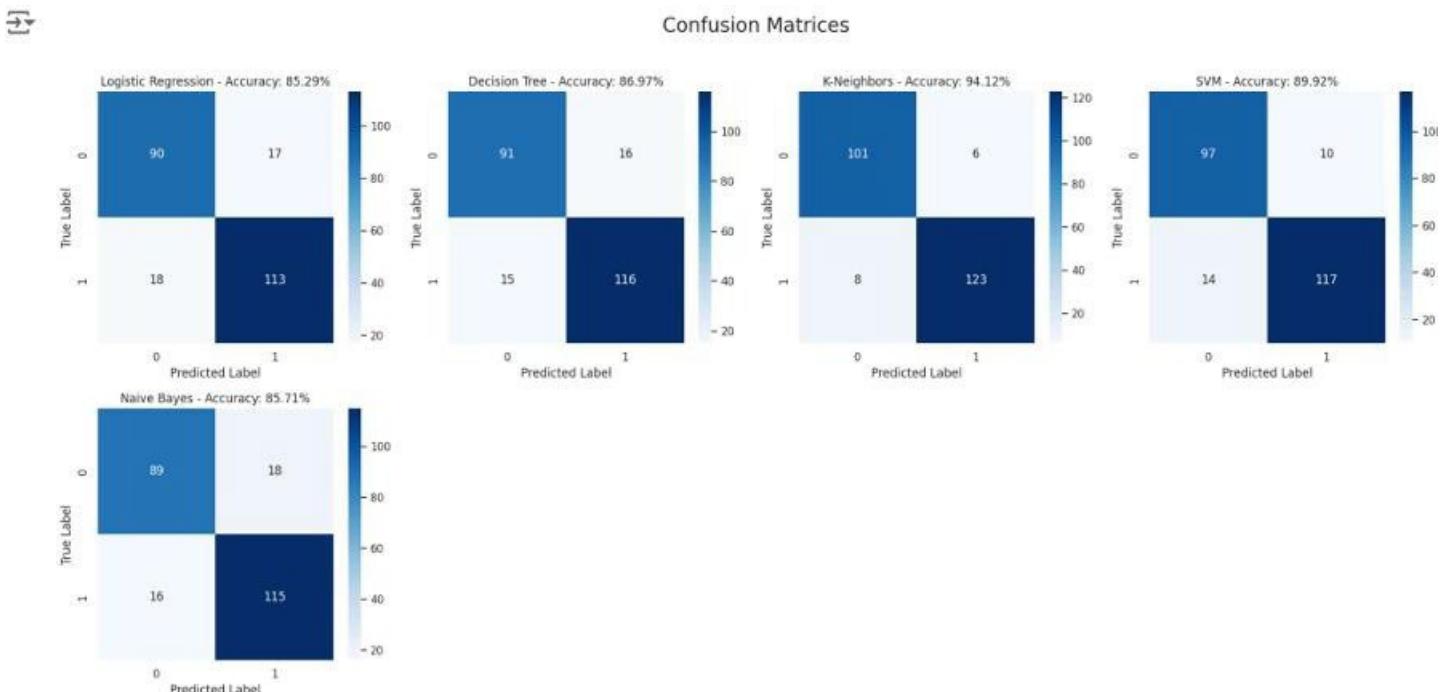
for (name, cm, acc_score), ax in zip(classifiers, axes.flatten()):
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", ax=ax)

    ax.set_title(f"{name} - Accuracy: {acc_score * 100:.2f}%")
    ax.set_xlabel("Predicted Label")
    ax.set_ylabel("True Label")

for i in range(num_classifiers, num_rows * num_cols):
    fig.delaxes(axes.flatten()[i])

plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()

```



4. Classification Report Comparison Bar Plots

```
# Get classification reports as dictionaries
lr_cr = classification_report(Y_test, logreg_predict, output_dict=True)
dt_cr = classification_report(Y_test, DT_predict, output_dict=True)
knn_cr = classification_report(Y_test, KNN_predict, output_dict=True)
svm_cr = classification_report(Y_test, SVM_predict, output_dict=True)
nb_cr = classification_report(Y_test, NB_predict, output_dict=True)

classification_reports = [lr_cr, dt_cr, knn_cr, svm_cr, nb_cr]

f1_scores = {}
recall_scores = {}
precision_scores = {}

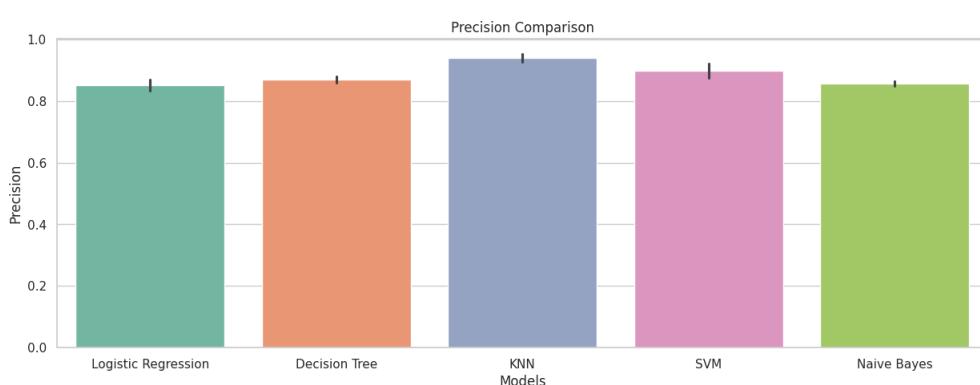
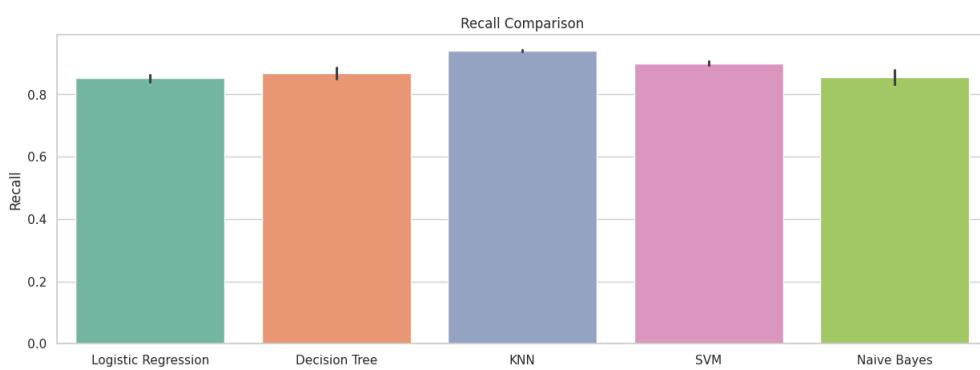
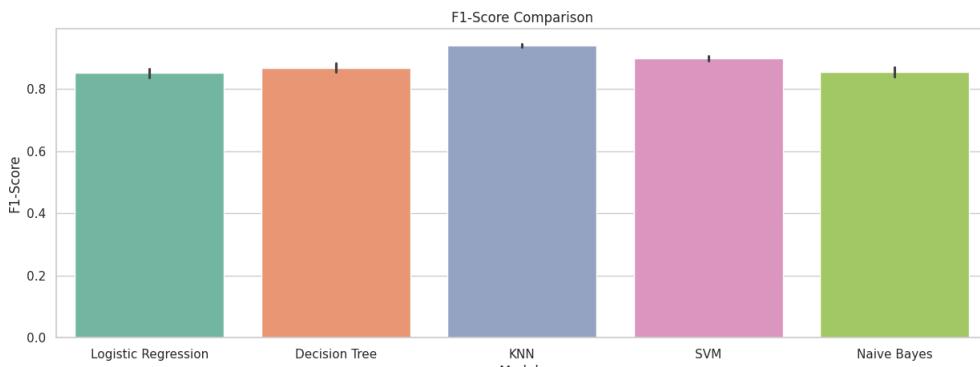
# Store f1-score, recall, and precision scores in lists
for name, cr in zip(
    [
        "Logistic Regression",
        "Decision Tree",
        "KNN",
        "SVM",
        "Naive Bayes",
    ],
    classification_reports,
):
    f1_scores[name] = [
        cr[label]["f1-score"] for label in cr.keys() if label.isnumeric()
    ]
    recall_scores[name] = [
        cr[label]["recall"] for label in cr.keys() if label.isnumeric()
    ]
    precision_scores[name] = [
        cr[label]["precision"] for label in cr.keys() if label.isnumeric()
    ]

# Create pandas dataframes from the lists
df_f1 = pd.DataFrame(
    f1_scores, index=[str(i) for i in range(1, len(f1_scores["Decision Tree"]) + 1)]
)
df_recall = pd.DataFrame(
    recall_scores,
    index=[str(i) for i in range(1, len(recall_scores["Decision Tree"]) + 1)],
)
df_precision = pd.DataFrame(
    precision_scores,
    index=[str(i) for i in range(1, len(precision_scores["Decision Tree"]) + 1)],
)

# Plot accuracy comparison
plt.figure(figsize=(15, 5))
sns.barplot(data=df_f1, palette="Set2")
plt.title("F1-Score Comparison")
plt.xlabel("Models")
plt.ylabel("F1-Score")
plt.show()
```

```
# Plot recall comparison
plt.figure(figsize=(15, 5))
sns.barplot(data=df_recall, palette="Set2")
plt.title("Recall Comparison")
plt.xlabel("Models")
plt.ylabel("Recall")
plt.show()
```

```
# Plot precision comparison
plt.figure(figsize=(15, 5))
sns.barplot(data=df_precision, palette="Set2")
plt.title("Precision Comparison")
plt.xlabel("Models")
plt.ylabel("Precision")
plt.show()
```



Conclusion

By comparing various Supervised Learning Algorithm on Heart Disease Dataset. We can say that KNN is the best Supervised Learning Algorithm for these dataset.

Experiment - 11

Aim: Compare the various Unsupervised learning algorithm by using the appropriate datasets.

Comparison of various Unsupervised Learning Algorithms:

```
from sklearn import datasets
from sklearn.cluster import KMeans
from matplotlib import pyplot as plt
# Load the Dataset
ds = datasets.load_wine()
x = ds.data # independent variable
x.shape

→ (178, 13)

import pandas as pd
x_df = pd.DataFrame(x, columns=ds.feature_names)
x_df.head()

→
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenols	proanthocyanins	color_i
0	14.23	1.71	2.43	15.6	127.0	2.80	3.06	0.28	2.29	
1	13.20	1.78	2.14	11.2	100.0	2.65	2.76	0.26	1.28	
2	13.16	2.36	2.67	18.6	101.0	2.80	3.24	0.30	2.81	
3	14.37	1.95	2.50	16.8	113.0	3.85	3.49	0.24	2.18	
4	13.24	2.59	2.87	21.0	118.0	2.80	2.69	0.39	1.82	

Next steps: [Generate code with x_df](#) [View recommended plots](#) [New interactive sheet](#)

```
x_df.info()

→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 178 entries, 0 to 177
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   alcohol          178 non-null    float64
 1   malic_acid       178 non-null    float64
 2   ash               178 non-null    float64
 3   alcalinity_of_ash 178 non-null    float64
 4   magnesium        178 non-null    float64
 5   total_phenols    178 non-null    float64
 6   flavanoids        178 non-null    float64
 7   nonflavanoid_phenols 178 non-null    float64
 8   proanthocyanins  178 non-null    float64
 9   color_intensity   178 non-null    float64
 10  hue               178 non-null    float64
 11  od280/od315_of_diluted_wines 178 non-null    float64
 12  proline          178 non-null    float64
dtypes: float64(13)
memory usage: 18.2 KB
```

```
x_df.describe()
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenols	proanthocy.
count	178.000000	178.000000	178.000000	178.000000	178.000000	178.000000	178.000000	178.000000	178.000000
mean	13.000618	2.336348	2.366517	19.494944	99.741573	2.295112	2.029270	0.361854	1.51
std	0.811827	1.117146	0.274344	3.339564	14.282484	0.625851	0.998859	0.124453	0.51
min	11.030000	0.740000	1.360000	10.600000	70.000000	0.980000	0.340000	0.130000	0.4
25%	12.362500	1.602500	2.210000	17.200000	88.000000	1.742500	1.205000	0.270000	1.21
50%	13.050000	1.865000	2.360000	19.500000	98.000000	2.355000	2.135000	0.340000	1.51
75%	13.677500	3.082500	2.557500	21.500000	107.000000	2.800000	2.875000	0.437500	1.91
max	14.830000	5.800000	3.230000	30.000000	162.000000	3.880000	5.080000	0.660000	3.51

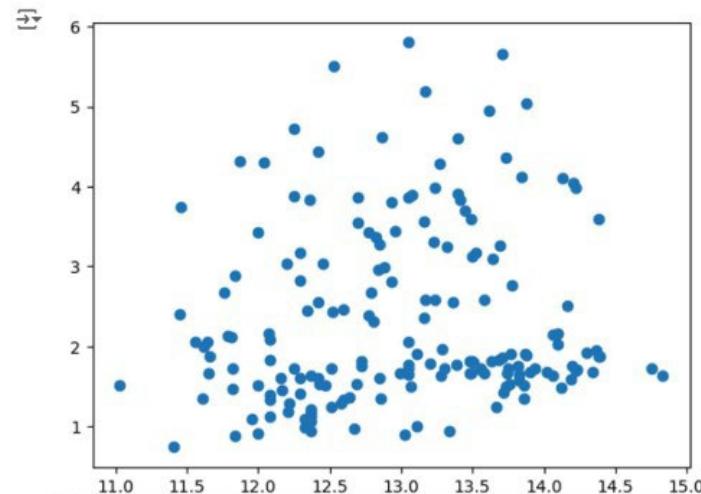
```
x_df.isnull().sum()
```

	0
alcohol	0
malic_acid	0
ash	0
alcalinity_of_ash	0
magnesium	0
total_phenols	0
flavanoids	0
nonflavanoid_phenols	0
proanthocyanins	0
color_intensity	0
hue	0
od280/od315_of_diluted_wines	0
proline	0

```
x_df.duplicated().sum()
```

```
→ 0
```

```
plt.scatter(x[:,0], x[:,1]) # (x[:,0], x[:,1]) = (x[all rows, 1st column], x[all rows, 2nd column])
plt.show()
```

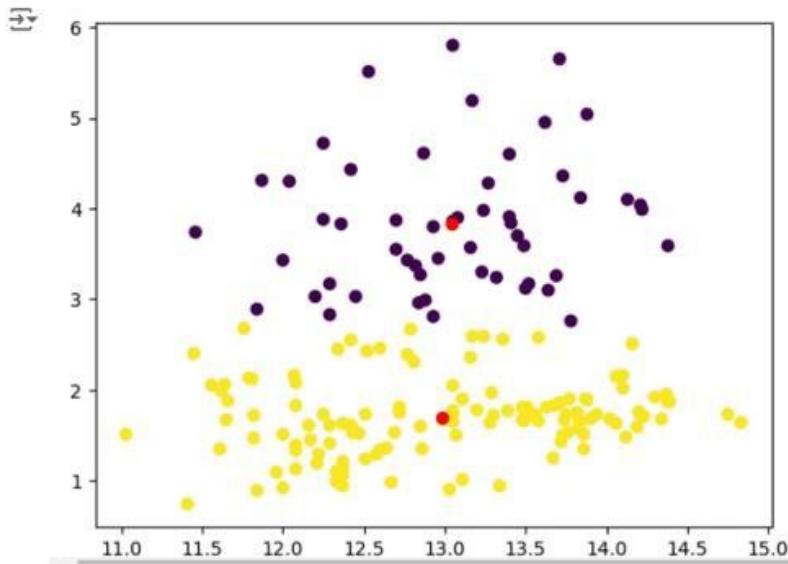


```
cluster_2_model = KMeans(n_clusters=2)
cluster_labels = cluster_2_model.fit_predict(x[:,[0,1]])

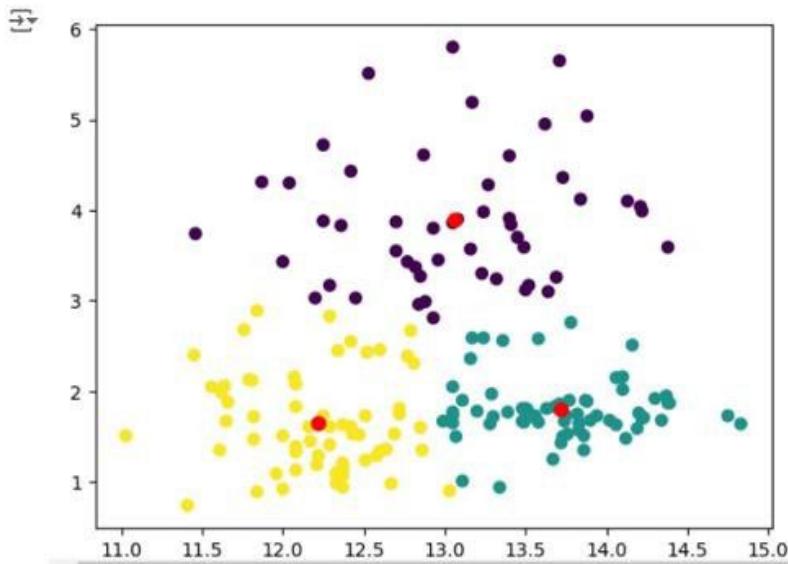
cluster_labels[:10], cluster_labels[-10:]

→ (array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int32),
 array([0, 1, 1, 0, 0, 1, 1, 1, 0, 1], dtype=int32))
```

```
cluster_2_model = KMeans(n_clusters=2)
cluster_labels = cluster_2_model.fit_predict(x[:,[0,1]])
plt.scatter(x[:,0], x[:,1], c=cluster_labels)
centroids = cluster_2_model.cluster_centers_
plt.scatter(centroids[:,0], centroids[:,1], c='red')
plt.show()
```



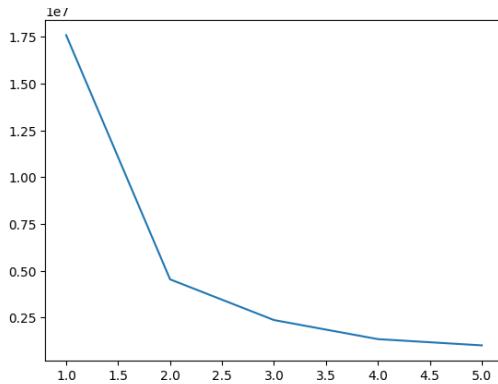
```
cluster_3_model = KMeans(n_clusters=3)
cluster_labels = cluster_3_model.fit_predict(x[:,[0,1]])
plt.scatter(x[:,0], x[:,1], c=cluster_labels)
centroids = cluster_3_model.cluster_centers_
plt.scatter(centroids[:,0], centroids[:,1], s=50, color='red')
plt.show()
```



```
# Using Elbow method to find best values of k
k_values = [1, 2, 3, 4, 5]
wcss_values = []
```

```
for k_value in k_values:
    model = KMeans(n_clusters=k_value)
    model.fit(x)
    wcss_value = model.inertia_
    wcss_values.append(wcss_value)
```

```
plt.plot(k_values, wcss_values)
plt.show()
```



```
from sklearn.metrics import silhouette_score
```

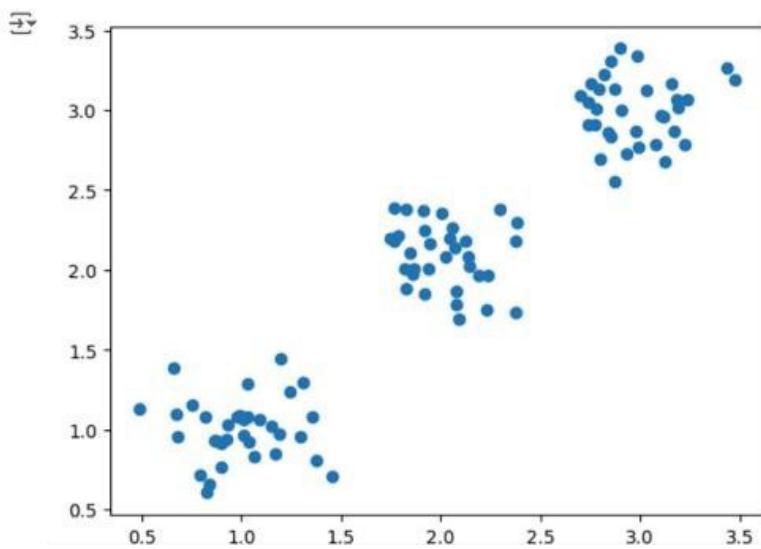
```
k_values = [2,3,4,5]
wcss_values = []

for k_value in k_values:
    model = KMeans(n_clusters=k_value)
    cluster_labels = model.fit_predict(x[:,[0, 1]])
    model_score = silhouette_score(x[:,[0,1]], cluster_labels)

    print(f"Cluster score: {model_score} for {k_value}")
```

```
Cluster score: 0.47745578237294 for 2
Cluster score: 0.4806665140821557 for 3
Cluster score: 0.45020918927738124 for 4
Cluster score: 0.37104988795151717 for 5
```

```
x_data = datasets.make_blobs(n_samples = 100, centers=[[1, 1], [2, 2], [3, 3]], cluster_std = 0.2, random_state = 0)
plt.scatter(x_data[0][:,0], x_data[0][:,1])
plt.show()
```



Experiment - 12

Aim: Build an Artificial Neural Network by implementing the Backpropagation algorithm and test the same using appropriate data sets.

Backpropagation Algorithm:

```

import tensorflow as tf
from tensorflow.keras import datasets

(x_train, y_train), (x_test, y_test) = datasets.boston_housing.load_data()

print('Training set: ', x_train.shape, y_train.shape)
print('Test set: ', x_test.shape, y_test.shape)

→ Training set: (404, 13) (404,)
    Test set: (102, 13) (102,)

y_train[:5]

→ array([15.2, 42.3, 50. , 21.1, 17.7])

x_train[0]

→ array([ 1.23247,   0.        ,   8.14     ,   0.        ,
       91.7      ,   3.9769   ,   4.        ,  307.      ,
      18.72     ])
       ,   0.538     ,   6.142     ,
       21.        ,  396.9      ,
       ,   18.72     ])

# Build ANN Arch for Regression

from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense, Input
from tensorflow.keras.utils import plot_model

model_reg = Sequential([
    Input((13,)),
    Dense(1, activation='linear')
])

# model_reg.summary()
# plot_model(model_reg, show_shapes=True, show_layer_activations=True)

# Model Training
model_reg.compile(optimizer='sgd', loss='msle', metrics=['accuracy', 'msle'])
model_reg.fit(x_train, y_train, epochs=10)

→ Epoch 1/10
13/13 ━━━━━━━━ 1s 2ms/step - accuracy: 0.0000e+00 - loss: 6.4965 - msle:
Epoch 2/10
13/13 ━━━━━━ 0s 1ms/step - accuracy: 0.0000e+00 - loss: 9.5771 - msle:
Epoch 3/10
13/13 ━━━━ 0s 1ms/step - accuracy: 0.0000e+00 - loss: 9.7181 - msle:
Epoch 4/10
13/13 ━━ 0s 1ms/step - accuracy: 0.0000e+00 - loss: 9.5237 - msle:
Epoch 5/10
13/13 ━ 0s 2ms/step - accuracy: 0.0000e+00 - loss: 9.4521 - msle:
Epoch 6/10
13/13 0s 1ms/step - accuracy: 0.0000e+00 - loss: 9.6827 - msle:
Epoch 7/10

```

```
13/13 ━━━━━━━━ 0s 2ms/step - accuracy: 0.0000e+00 - loss: 9.5076 - msle:  
Epoch 8/10  
13/13 ━━━━━━ 0s 1ms/step - accuracy: 0.0000e+00 - loss: 9.6568 - msle:  
Epoch 9/10  
13/13 ━━━━ 0s 2ms/step - accuracy: 0.0000e+00 - loss: 9.7762 - msle:  
Epoch 10/10  
13/13 ━━ 0s 2ms/step - accuracy: 0.0000e+00 - loss: 9.7529 - msle:  
<keras.src.callbacks.history.History at 0x7d5dd8624940>
```

```
# Evaluate the model  
model_reg.evaluate(x_test, y_test)
```

```
→ 4/4 ━━━━ 0s 3ms/step - accuracy: 0.0000e+00 - loss: 9.6913 - msle: 9.  
[9.806747436523438, 0.0, 9.806747436523438]
```



```
# Build ANN Arch for Regression
```

```
from tensorflow.keras import Sequential  
from tensorflow.keras.layers import Dense, Input  
from tensorflow.keras.utils import plot_model
```

```
model_reg_1 = Sequential()  
model_reg_1.add(Dense(50, activation='relu', input_shape=(x_train.shape[1], )))  
model_reg_1.add(Dense(25, activation='relu'))  
model_reg_1.add(Dense(12, activation='relu'))  
model_reg_1.add(Dense(1, activation='linear'))
```

```
# model_reg_1.summary()  
# plot_model(model_reg_1, show_shapes=True, show_layer_activations=True)
```

```
# Model Training
```

```
model_reg_1.compile(optimizer='adam', loss='msle', metrics=['accuracy', 'msle'])  
model_reg_1.fit(x_train, y_train, epochs=10, validation_data=(x_test, y_test), batch_size
```

```
→ Epoch 1/10  
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning  
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)  
7/7 ━━━━ 2s 41ms/step - accuracy: 0.0000e+00 - loss: 0.3993 - msle: 0  
Epoch 2/10  
7/7 ━━━━ 0s 6ms/step - accuracy: 0.0000e+00 - loss: 0.1851 - msle: 0.  
Epoch 3/10  
7/7 ━━━━ 0s 10ms/step - accuracy: 0.0000e+00 - loss: 0.1372 - msle: 0  
Epoch 4/10  
7/7 ━━━━ 0s 9ms/step - accuracy: 0.0000e+00 - loss: 0.1228 - msle: 0.  
Epoch 5/10  
7/7 ━━━━ 0s 6ms/step - accuracy: 0.0000e+00 - loss: 0.1111 - msle: 0.  
Epoch 6/10  
7/7 ━━━━ 0s 9ms/step - accuracy: 0.0000e+00 - loss: 0.0861 - msle: 0.  
Epoch 7/10  
7/7 ━━━━ 0s 9ms/step - accuracy: 0.0000e+00 - loss: 0.0781 - msle: 0.  
Epoch 8/10  
7/7 ━━━━ 0s 9ms/step - accuracy: 0.0000e+00 - loss: 0.0768 - msle: 0.  
Epoch 9/10  
7/7 ━━━━ 0s 11ms/step - accuracy: 0.0000e+00 - loss: 0.0980 - msle: 0  
Epoch 10/10
```

7/7 0s 6ms/step - accuracy: 0.0000e+00 - loss: 0.0816 - msle: 0.
<keras.src.callbacks.history.History at 0x7d5dd844fdc0>

```
# ANN for binary classification
```

```
import pandas as pd
from sklearn.preprocessing import LabelEncoder
```

```
bank_ds = pd.read_csv('/content/Bank(1 ).csv', delimiter=';')
```

```
le = LabelEncoder()
```

```
for col in ['job', 'marital', 'default', 'education', 'housing', 'loan', 'contact', 'mont
    bank_ds[col] = le.fit_transform(bank_ds[col])
bank_ds.head()
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month
0	30	10	1	0	0	1787	0	0	0	19	10
1	33	7	1	1	0	4789	1	1	0	11	8
2	35	4	2	2	0	1350	1	0	0	16	0
3	30	4	1	2	0	1476	1	1	2	3	6
4	59	1	1	1	0	0	1	0	2	5	8

Next
steps:

Generate code
with bank_ds



View recommended
plots

New interactive
sheet

```
x = bank_ds.iloc[:, :-1].values
y = bank_ds.iloc[:, -1].values
```

```
# import numpy as np
```

```
# x.shape
```

```
# for i in range(x.shape[1]):
#     x[:,i] = x[:,i]/np.max(x[:,i])
```

```
from sklearn.preprocessing import StandardScaler
```

```
x_std = StandardScaler().fit_transform(x)
```

```
x_std[0]
```

```
array([-1.05626965,  1.71680374, -0.24642938, -1.64475535, -0.1307588 ,
       0.12107186, -1.14205138, -0.42475611, -0.72364152,  0.37405206,
      1.48541444, -0.7118608 , -0.57682947, -0.4072183 , -0.32041282,
      0.44441328])
```

```
# ANN

model_bc = Sequential()
model_bc.add(Dense(20, activation='relu', input_shape=(16,)))
model_bc.add(Dense(5, activation='relu'))
model_bc.add(Dense(1, activation='sigmoid')) #binary classification take sigmoid as activ

model_bc.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model_bc.fit(x, y, epochs=10, batch_size=128, validation_split=0.2)

→ Epoch 1/10
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
29/29 ━━━━━━━━ 1s 9ms/step - accuracy: 0.5664 - loss: 6.5969 - val_accuracy: 0.5664
Epoch 2/10
29/29 ━━━━━━━━ 0s 3ms/step - accuracy: 0.6788 - loss: 3.1804 - val_accuracy: 0.6788
Epoch 3/10
29/29 ━━━━━━━━ 0s 3ms/step - accuracy: 0.7214 - loss: 2.1814 - val_accuracy: 0.7214
Epoch 4/10
29/29 ━━━━━━━━ 0s 2ms/step - accuracy: 0.7086 - loss: 1.8458 - val_accuracy: 0.7086
Epoch 5/10
29/29 ━━━━━━━━ 0s 3ms/step - accuracy: 0.8008 - loss: 1.3134 - val_accuracy: 0.8008
Epoch 6/10
29/29 ━━━━━━━━ 0s 4ms/step - accuracy: 0.8270 - loss: 1.0517 - val_accuracy: 0.8270
Epoch 7/10
29/29 ━━━━━━━━ 0s 3ms/step - accuracy: 0.8126 - loss: 1.0897 - val_accuracy: 0.8126
Epoch 8/10
29/29 ━━━━━━━━ 0s 3ms/step - accuracy: 0.8322 - loss: 0.8262 - val_accuracy: 0.8322
Epoch 9/10
29/29 ━━━━━━━━ 0s 3ms/step - accuracy: 0.8456 - loss: 0.7307 - val_accuracy: 0.8456
Epoch 10/10
29/29 ━━━━━━━━ 0s 3ms/step - accuracy: 0.8547 - loss: 0.7296 - val_accuracy: 0.8547
<keras.src.callbacks.history.History at 0x7d5dc4a718a0>
```

```
# prediction of new
pred_proba = model_bc.predict(x)

pred = []

for proba in pred_proba:
    if proba>=0.5:
        pred.append(1)
    else:
        pred.append(0)

→ 142/142 ━━━━━━━━ 0s 2ms/step

from sklearn.metrics import classification_report

cls = classification_report(y, pred)
print(cls)

→
      precision    recall   f1-score   support
      0          0.91     0.94     0.92     1000
```

1	0.38	0.29	0.33	521
accuracy			0.86	4521
macro avg	0.64	0.61	0.63	4521
weighted avg	0.85	0.86	0.86	4521

```
from tensorflow.keras import losses
from tensorflow.keras import optimizers

# ANN for multiclass class

from sklearn import datasets
from tensorflow.keras.utils import to_categorical

iris_ds = datasets.load_iris()

x = iris_ds.data
y = iris_ds.target

y_cat = to_categorical(y, num_classes=3) #one-hot encoding

model_mc = Sequential()
model_mc.add(Dense(10, activation='relu', input_shape=(4,)))
model_mc.add(Dense(5, activation='relu'))
model_mc.add(Dense(3, activation='softmax')) #multiclass classification take softmax

model_mc.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model_mc.fit(x, y_cat, epochs=4, batch_size=64, validation_split=0.2)

→ Epoch 1/4
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning:
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
2/2 ━━━━━━━━ 1s 219ms/step - accuracy: 0.4236 - loss: 1.1951 - val_accuracy: 0.4236 - val_loss: 1.1951
Epoch 2/4
2/2 ━━━━━━━━ 0s 33ms/step - accuracy: 0.4132 - loss: 1.1791 - val_accuracy: 0.4132 - val_loss: 1.1791
Epoch 3/4
2/2 ━━━━━━━━ 0s 28ms/step - accuracy: 0.3924 - loss: 1.1723 - val_accuracy: 0.3924 - val_loss: 1.1723
Epoch 4/4
2/2 ━━━━━━━━ 0s 27ms/step - accuracy: 0.4080 - loss: 1.1435 - val_accuracy: 0.4080 - val_loss: 1.1435
<keras.src.callbacks.history.History at 0x7d5dc99d5870>
```

```
import numpy as np

x.shape
np.unique(y)

→ array([0, 1, 2])

y_pred_proba = model_mc.predict(x)

y_pred = []
for proba in y_pred_proba:
```

```
y_pred.append(np.argmax(proba))

[?] 5/5 ━━━━━━━━ 0s 10ms/step

print(classification_report(y, y_pred))

precision    recall   f1-score   support
          0       0.00     0.00     0.00      50
          1       0.33     1.00     0.50      50
          2       0.00     0.00     0.00      50

accuracy                           0.33      150
macro avg       0.11     0.33     0.17      150
weighted avg    0.11     0.33     0.17      150
```